

Operating Systems Design

6. Synchronization

Paul Krzyzanowski
pxk@cs.rutgers.edu

Concurrency

Concurrent threads/processes (informal)

- Two processes are concurrent if they run at the same time or if their execution is interleaved *in any order*

Asynchronous

- The processes require occasional synchronization

Independent

- They do not have any reliance on each other

Synchronous

- Frequent synchronization with each other – order of execution is guaranteed

Parallel

- Processes run at the same time on separate processors

Race Conditions

A **race condition** is a bug:

- The outcome of concurrent threads are unexpectedly dependent on a specific sequence of events.

Example

- Your current bank balance is \$1,000.
- Withdraw \$500 from an ATM machine while a \$5,000 direct deposit is coming in

Withdrawal

- Read account balance
- Subtract 500
- Write account balance

Deposit

- Read account balance
- Add 5000
- Write account balance

Possible outcomes:

Total balance = \$5500, \$500, \$6000

Synchronization

Synchronization deals with developing techniques to avoid race conditions

Something as simple as

$x = x + 1;$

May have have a race condition:

`movl _x (%rip), %eax`

`addl $1, %eax`

`movl %eax, _x (%rip)`

←
← } Potential points of
preemption for a race
condition

Mutual Exclusion

Critical section:

- Region in a program where race conditions can arise

Mutual exclusion:

- Allow only one thread to access a critical section at a time

Deadlock:

- A thread is perpetually blocked (circular dependency on resources)

Starvation:

- A thread is perpetually denied resources

Livelock:

- Threads run but no progress in execution

Controlling critical section access: locks

- Grab and release locks around critical sections
- Wait if you cannot get a lock

Withdrawal

Enter Critical
Section

- Acquire(`transfer_lock`)

Critical
Section

- Read account balance
- Subtract 500
- Write account balance

Exit Critical
Section

- Release(`transfer_lock`)

Deposit

- Acquire(`transfer_lock`)

Enter Critical
Section

- Read account balance
- Add 5000
- Write account balance

Critical
Section

- Release(`transfer_lock`)

Exit Critical
Section

The Critical Section Problem

Design a protocol to allow threads to enter a critical section

Conditions for a solution

- **Mutual exclusion**: No threads may be inside the same critical sections simultaneously
- **Progress**: If no thread is executing in its critical section but one or more threads want to enter, the selection of a thread cannot be delayed indefinitely.
 - If one thread wants to enter, it should be permitted to enter.
 - If multiple threads want to enter, exactly one should be selected.
- **Bounded waiting**: No thread should wait forever to enter a critical section
- No thread running outside its critical section may block others
- A good solution will make no assumptions on:
 - No assumptions on # processors
 - No assumption on # threads/processes
 - Relative speed of each thread

Critical sections & the kernel

- Multiprocessors
 - Multiple processes on different processors may access the kernel simultaneously
 - Interrupts may occur on multiple processors simultaneously
- Preemptive kernels
 - **Preemptive kernel**: process can be preempted while running in kernel mode
 - **Nonpreemptive kernel**: processes running in kernel mode cannot be preempted (but interrupts can still occur!)
- Single processor, nonpreemptive kernel: free from race conditions

Solution #1: Disable Interrupts

Disable all system interrupts before entering a critical section and re-enable them when leaving

Bad!

- Gives the thread too much control over the system
- Stops time updates and scheduling
- What if the logic in the critical section goes wrong?
- What if the critical section has a dependency on some other interrupt, thread, or system call?
- What about multiple processors? Disabling interrupts affects just one processor

Advantage

- Simple, guaranteed to work
- Was often used in the uniprocessor kernels

Solution #2: Software Test & Set Locks

Keep a shared lock variable:

```
while (locked) ;  
locked = 1;  
/* do critical section */  
locked = 0;
```

Disadvantage:

- Buggy! There's a race condition in setting the lock

Advantage:

- Simple to understand. It's been used for things such as locking mailbox files

Solution #3: Lockstep Synchronization

Take turns

Thread 0

```
while (turn != 0);  
critical_section();  
turn = 1;
```

Thread 1

```
while (turn != 1);  
critical_section();  
turn = 0;
```

Disadvantages:

- Tight loop that spins waiting for a turn: **busy waiting** or **spin lock**
- Forces strict alternation; if thread 2 is really slow, thread 1 is slowed down with it

Software solutions for mutual exclusion

- Peterson's solution (page 221 of text)
- Others
- Disadvantages:
 - Difficult to implement correctly – have to rely on `volatile` data types to ensure that compilers don't make the wrong optimizations
 - Relies on busy waiting

Let's turn to hardware for help

Help from the processor

Atomic (indivisible) CPU instructions that help us get locks

- Test-and-set
- Compare-and-swap
- Fetch-and-Increment

Test & Set

Test-and-set

```
ATOMIC { int test_and_set(int *x) {  
        last_value = *x;  
        *x = 1;  
        return last_value;  
    }  
}
```

Set the lock but get told if it already was set (in which case you don't have it)

```
while (test_and_set(&lock)) ;  
/* do critical section */  
lock = 0;
```


Compare & swap (CAS)

Compare the value of a memory location with an old value.
If they match then replace with a new value

ATOMIC {

```
int compare_and_swap(int *x, int old, int new) {  
    int save = *x;  
    if (save == old)  
        *x = new;  
    return save; /* always return location contents */  
}
```

Avoid the race condition.

Set *locked* to 1 only if *locked* is still set to 0.

```
while (compare_and_swap(&locked, 0, 1) != 0) ;  
    /* spin until locked == 0 */  
/* if we got here, locked got set to 1 and we have it */  
/* do critical section */  
locked = 0; /* release the lock */
```

Fetch & Increment

Increment a memory location; return previous value

```
ATOMIC { int fetch_and_increment(int *x) {  
        last_value = *x;  
        *x = *x + 1;  
        return last_value;  
    }  
}
```

Fetch & Increment

Check that it's your turn for the critical section

```
ticket = 0; turn = 0;  
...  
myturn = fetch_and_increment(&ticket);  
while (turn != myturn) ;  
/* do critical section */  
fetch_and_increment(&turn);
```



turn



ticket

Spin locks

- All these techniques rely on **spin locks**
- Wastes CPU cycles
- The process with the lock may not be allowed to run!
 - Lower priority process obtained a lock
 - Higher priority process is always ready to run but loops on trying to get the lock
 - Scheduler always schedules the higher-priority process
 - **Priority inversion**
 - If the low priority process would get to run & release its lock, it would then accelerate the time for the high priority process to get a chance to get the lock and do useful work
 - Try explaining that to a scheduler!

Priority Inheritance

- Technique to avoid priority inversion
- Increase the priority of any process to the maximum of any process waiting on any resource for which the process has a lock
- When the lock is released, the priority goes to its normal level

Spin locks aren't great

Can we block until we can get the critical section?

How about this?

```
public class Lock
{
    private int val = UNLOCKED;
    private ThreadQueue waitQueue = new ThreadQueue();

    public void acquire() {
        Thread me = Thread.currentThread();
        while (TestAndSet(val) == LOCKED) {
            waitQueue.waitForAccess(me); // Put self in queue
            Thread.sleep();              // Put self to sleep
        }
        // Got the lock
    }
    public void release() {
        Thread next = waitQueue.nextThread();
        val = UNLOCKED;
        if (next != null)
            next.ready(); // Wake up a waiting thread
    }
}
```

Sorry...

- Accessing the wait queue is a critical section
 - Need to add mutual exclusion
- Need extra lock check in *acquire*
 - Thread may find the lock busy
 - Another thread may release the lock but before the first thread enqueues itself

Semaphores

- Count # of wake-ups saved for future use
- Two atomic operations:

```
down(sem s) {  
    if (s > 0)  
        s = s - 1;  
    else  
        sleep on event s  
}
```

```
up(sem s) {  
    if (someone is waiting on s)  
        wake up one of the threads  
    else  
        s = s + 1;  
}
```

```
//initialize  
mutex = 1;
```

```
down(&mutex)
```

```
// critical section
```

```
up(&mutex)
```

Binary semaphore

Semaphores

Count the number of threads that may enter a critical section at any given time.

- Each *down* decreases the number of future accesses
- When no more are allowed, processes have to wait
- Each *up* lets a waiting process get in

Producer-Consumer example

- Producer
 - Generates items that go into a buffer
 - Maximum buffer capacity = N
 - If the producer fills the buffer, it must wait (sleep)
- Consumer
 - Consumes things from the buffer
 - If there's nothing in the buffer, it must wait (sleep)
- This is also known as the *Bounded-Buffer Problem*

Producer-Consumer example

```
sem mutex=1, empty=N, full=0;
producer() {
    for (;;) {
        produce_item(&item);    /* produce something */
        down(&empty);           /* decrement empty count */
        down(&mutex);           /* start critical section */
        enter_item(item);       /* put item in buffer */
        up(&mutex);             /* end critical section */
        up(&full);              /* +1 full slot */
    }
}
consumer() {
    for (;;) {
        down(&full);            /* one less item */
        down(&mutex);           /* start critical section */
        remove_item(item);     /* get the item from the buffer */
        up(&mutex);             /* end critical section */
        up(&empty);            /* one more empty slot */
        consume_item(item);    /* consume it */
    }
}
```

Readers-Writers example

- Shared data store (e.g., database)
- Multiple processes can read concurrently
- Only one process can write at a time
 - And no readers can read while the writer is writing

Readers-Writers example

```
sem mutex=1; /* critical sections used only by the reader */
sem canwrite=1; /* critical section for N readers vs. 1 writer */
int readcount = 0; /* number of concurrent readers */
```

```
writer() {
    for (;;) {
        down(&canwrite);          /* block if we cannot write */
        // write data
        up(&canwrite);            /* end critical section */
    }
}
```

Readers-Writers example

```
sem mutex=1; /* critical sections used only by the reader */
sem canwrite=1; /* critical section for N readers vs. 1 writer */
int readcount = 0; /* number of concurrent readers */
```

```
reader() {
critical section {
    for (;;) {
        down(&mutex);
        readcount++;
        if (readcount == 1)
            down(canwrite); /* sleep or disallow the writer from writing */
        up(&mutex);
        // do the read
        down(&mutex);
        readcount--;
        if (readcount == 0)
            up(writer); /* no more readers! Allow the writer access */
        up(&mutex);
        // other stuff
    }
}
```

Event Counters

Avoid race conditions without using mutual exclusion

Three operations:

- **read**(E): return the current value of event counter E
- **advance**(E): increment E (atomically)
- **await**(E, v): wait until E has a value $\geq v$

Producer-Consumer example

```
#define N 4          /* four slots in the buffer */
event_counter in=0;  /* number of items inserted into buffer */
event_counter out=0; /* number of items removed from buffer */

producer() {
    int item, sequence=0;
    for (;;) {
        produce_item(&item);    /* produce something */
        sequence++;             /* item # of item produced */
        —————→ await(out, sequence-N); /* wait until there's room */ (0≥-3), (0≥-2), ...
        enter_item(item);       /* put item in buffer */
        —————→ advance(&in);          /* let consumer know there's one more item */
    }
}

consumer() {
    int item, sequence=0;
    for (;;) {
        sequence++;             /* item # we want to consume */
        —————→ await(in, sequence); /* wait until that item is present */ (0≥1)
        remove_item(item);      /* get the item from the buffer */
        —————→ advance(&out);        /* let producer know item's gone */
        consume_item(item);      /* consume it */
    }
}
```

Producer-Consumer example

```
#define N 4          /* four slots in the buffer */
event_counter in=0;  /* number of items inserted into buffer */
event_counter out=0; /* number of items removed from buffer */
producer() {
    int item, sequence=0;
    for (;;) {
        produce_item(&item); /* produce something */
        sequence++;          /* item # of item produced */
        —————→ await(out, sequence-N); /* wait until there's room */ (0≥-3), (0≥-2), ...
        enter_item(item);    /* put item in buffer */
        —————→ advance(&in); /* let consumer know there's one more item */
    }
}
```

Suppose the producer runs for a while and the consumer does not:

Iteration 1: sequence=1, out=0 *await*(0, 1-4): continue since $0 \geq -3$, in=1
Iteration 2: sequence=2, out=0 *await*(0, 2-4): continue since $0 \geq -2$, in=2
Iteration 3: sequence=3, out=0 *await*(0, 3-4): continue since $0 \geq -1$, in=3
Iteration 4: sequence=4, out=0 *await*(0, 4-4): continue since $0 \geq 0$, in=4
Iteration 5: sequence=5, out=0 *await*(0, 5-4): wait since $0 < 1$

Producer-Consumer example

```
#define N 4          /* four slots in the buffer */
event_counter in=0;  /* number of items inserted into buffer */
event_counter out=0; /* number of items removed from buffer */

consumer() {
    int item, sequence=0;
    for (;;) {
        sequence++;          /* item # we want to consume */
        → await(in, sequence); /* wait until that item is present */ (0 ≥ 1)
        remove_item(item);   /* get the item from the buffer */
        → advance(&out);     /* let producer know item's gone */
        consume_item(item);  /* consume it */
    }
}
```

Suppose the consumer runs first:

Iteration 1: sequence = 1, *await*(0, 1): sleep since $0 < 1$

When the producer runs its first iteration, it will increment *in*

The consumer's *await* will wake up since it's now *await*(1,1) and $1 \geq 1$

Condition Variables / Monitors

- Higher-level synchronization primitive
- Implemented by the programming language / APIs
- Two operations:
 - wait(*condition_variable*)
 - Block until *condition_variable* is “signaled”
 - signal(*condition_variable*)
 - Wake up one process that is waiting on the condition variable
 - Also called notify

Synchronization

Part II: Inter-Process Communication

Communicating processes

- Must:
 - Synchronize
 - Exchange data
- Message passing offers:
 - Data communication
 - Synchronization (via waiting for messages)
 - Works with processes on different machines

Message passing

- Two primitives:
 - send(destination, message)
 - receive(source, message)
- Operations may or may not be blocking

Producer-consumer example

```
#define N 4          /* number of slots in the buffer */

consumer() {
    int item, i;
    message m;

    for (i=0; i < N; ++i)
        send(producer, &m); /* send N empty messages */
    for (;;) {
        receive(producer, &m) /* get a message with the item */
        extract_item(&m, &item) /* take item out of message */
        send(producer, &m); /* send an empty reply */
        consume_item(item); /* consume it */
    }
}

producer() {
    int item;
    message m;


    for (;;) {
        produce_item(&item); /* produce something */
        receive(consumer, &m); /* wait for an empty message */
        build_message(&m, item); /* construct the message */
        send(consumer, &m); /* send it off */
    }
}
```


Messaging: Rendezvous

- Sending process blocked until receive occurs
- Receive blocks until a send occurs
- Advantages:
 - No need for message buffering if on same system
 - Easy & efficient to implement
 - Allows for tight synchronization
- Disadvantage:
 - Forces sender & receiver to run in lockstep

Messaging: Direct Addressing

- Sending process identifies receiving process
- Receiving process can identify sending process
 - Or can receive it as a parameter




Messaging: Indirect Addressing

- Messages sent to an intermediary data structure of FIFO queues
- Each queue is a mailbox
- Simplifies multiple readers


Mailboxes




Single sender, single reader



Multiple senders, single reader



Single sender, multiple readers



Multiple senders, multiple readers

Example: What you find in the world

IPC mechanisms in Windows

- Clipboard: central repository for sharing data among apps
- Dynamic Data Exchange (DDE) – [older] allows apps to exchange data in various formats; extension of clipboard
- COM: automatically start another app to access data via interfaces
- Data Copy: Windows messaging – two cooperating processes can copy data
- File Mapping: Memory mapped files
- Mailslots: one-way communication of short messages (including network)
- Anonymous pipes: for I/O redirection
- Named pipes
- RPC (remote procedure call)
- Sockets
- Semaphore objects

Linux

IPC Mechanisms in Linux

- Signals
- Pipes
- Named pipes (FIFOs)
- Semaphores
- Message queues: one or more writers & one or more readers
- Shared memory
- Memory-mapped files
- RPC (remote procedure calls)
- Sockets


Deadlocks

Four conditions must hold


1. Mutual exclusion
2. Hold and wait
3. Non-preemption of resources
 - Resources can only be released voluntarily
4. Circular wait

Deadlocks

- Resource allocation graph
 - Resource R_1 is allocated to process P_1 : **assignment edge**




- Resource R_1 is requested by process P_1 : **request edge**
:



- Deadlock is present when the graph has **cycles**

Deadlock example



Circular dependency among four processes and four resources

Dealing with deadlock

- **Deadlock prevention**
 - Ensure that at least one of the necessary conditions cannot hold
- **Deadlock avoidance**
 - Provide advance information to the OS on which resources a process will request.
 - OS can then decide if the process should wait
- **Ignore the problem**
 - Let the user deal with it (most common solution)

The End