

Operating Systems Design  
8. Memory Management

Paul Krzyzanowski  
pxk@cs.rutgers.edu

CPU Memory Access


CPU reads instructions and reads/write data from/to memory



Functional interface:  
`value = read(address)`  
`write(address, value)`

Monoprogramming


- Run one program at a time
- Share memory between the program and the OS



This was the model in old MS-DOS systems

Multiprogramming

- Keep more than one process in memory
- More processes in memory improves CPU utilization




- If a process spends 20% of its time computing, then would switching among 5 processes give us 100% CPU utilization?
- Not quite. For  $n$  processes, if  $p = \% \text{ time a process is blocked on I/O}$  then:  
 $\text{probability all are blocked} = p^n$
- CPU is not idle for  $(1-p^n)$  of the time
- 5 processes: 67% utilization

How do you access memory?

- **Absolute code**  
if you know where the program gets loaded (any relocation is done at link time)
- **Position independent code**  
all addresses are relative
- **Dynamically relocatable code**  
relocated at load time
- Or ... use **logical addresses**  
absolute code with addresses translated at runtime


Logical addressing

- **Memory management unit (MMU):**  
Real-time, on-demand translation between *logical* and *physical* addresses



### Relocatable addressing

- Base & limit
  - $\text{Physical address} = \text{logical address} + \text{base register}$
  - But first check that:  $\text{logical address} < \text{limit}$



### Multiple Fixed Partitions

- Divide memory into predefined partitions (segments)
  - Partitions don't have to be the same size
  - For example: a few big partitions and many small ones
- New process gets queued for a partition that can hold it
- Unused memory in a partition goes unused


### Variable partition multiprogramming

- Create partitions as needed
- New process gets queued
- OS tries to find a hole for it



### Variable partition multiprogramming

- Create partitions as needed
- New process gets queued
- OS tries to find a hole for it




### Variable partition multiprogramming

- What if a process needs more memory?
  - Always allocate some extra memory just in case
  - Find a hole big enough to relocate the process
- Combining holes
  - Memory compaction
  - Usually not done because of CPU time to move a lot of memory

### Segmentation hardware

- Divide a process into segments and place each segment into a partition of memory
  - Code segment, data segment, stack segment, etc.



### Allocation algorithms

- **First fit**: find the first hole that fits
- **Best fit**: find the hole that best fits the process
- **Worst fit**: find the largest available hole
  - *Why?* Maybe the remaining space will be big enough for another process. In practice, this algorithm does not work well.

### Paging

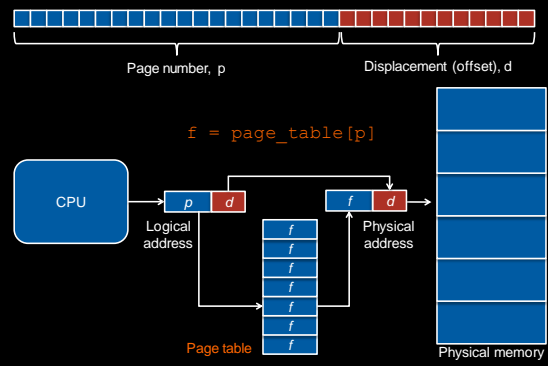
- Memory management scheme
  - Physical space can be non-contiguous
  - No fragmentation problems
  - No need for compaction
- Paging is implemented by the **Memory Management Unit (MMU)** in the processor

### Paging

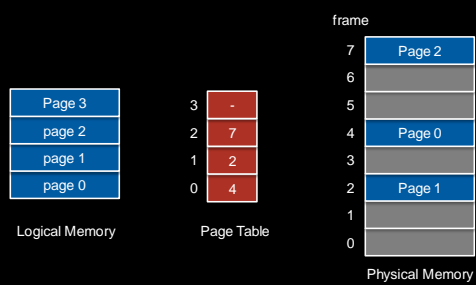
- Translation:
  - Divide physical memory into fixed-size blocks: **page frames**
  - A logical address is divided into blocks of the same size: **pages**
  - All memory accesses are translated: **page → page frame**
  - A page table maps pages to frames
- Example:
  - 32-bit address, 4 KB page size:
    - Top 20 bits identify the page number
    - Bottom 12 bits identify offset within the page/frame



### Page translation



### Logical vs. physical views of memory



### Hardware Implementation

- Where do you keep the page table? *In memory*
- Each process gets its own virtual address space
  - Each process has its own page table
  - Change the page table by changing a **page table base register**
- Memory translation is now slow!
  - To read a byte of memory, we need to read the page table first
  - Each memory access is now 2x slower!

## Hardware Implementation: TLB

- Cache frequently-accessed pages
  - Translation lookaside buffer (TLB)
  - Associative memory: key (page #) and value (frame #)
- TLB is on-chip & fast ... but small (64-1,024 entries)
- TLB miss: result not in the TLB
  - Need to do page table lookup in memory
- Hit ratio = % of lookups that come from the TLB

## Tagged TLB

- There is only one TLB per system
- When we context switch, we switch address spaces
  - New page table
  - TLB entries belong to the old address space
- Either:
  - Flush the TLB
  - Have a Tagged TLB:
    - Address space identifier (ASID)


## Protection

- An MMU can enforce memory protection
- Page table stores protection bits per frame
  - Valid/invalid: is there a frame mapped to this page?
  - Read-only
  - No execute
  - Dirty

## Multilevel (Hierarchical) page tables


- Most processes use only a small part of their address space
- Keeping an entire page table is wasteful
- E.g., 32-bit system with 4KB pages: 20-bit page table
  - 1,048,576 entries in a page table

## Multilevel page table



## Inverted page tables

- # of pages on a system may be huge
- # of page frames will be more manageable (limited by physical memory)
- Inverted page table
  - $i^{\text{th}}$  entry: contains info on what is in page frame  $i$
- Table access is no longer a simple index but a search
  - Use hashing and take advantage of associative memory



The End