

Operating Systems Design

10. Memory Management: Paging

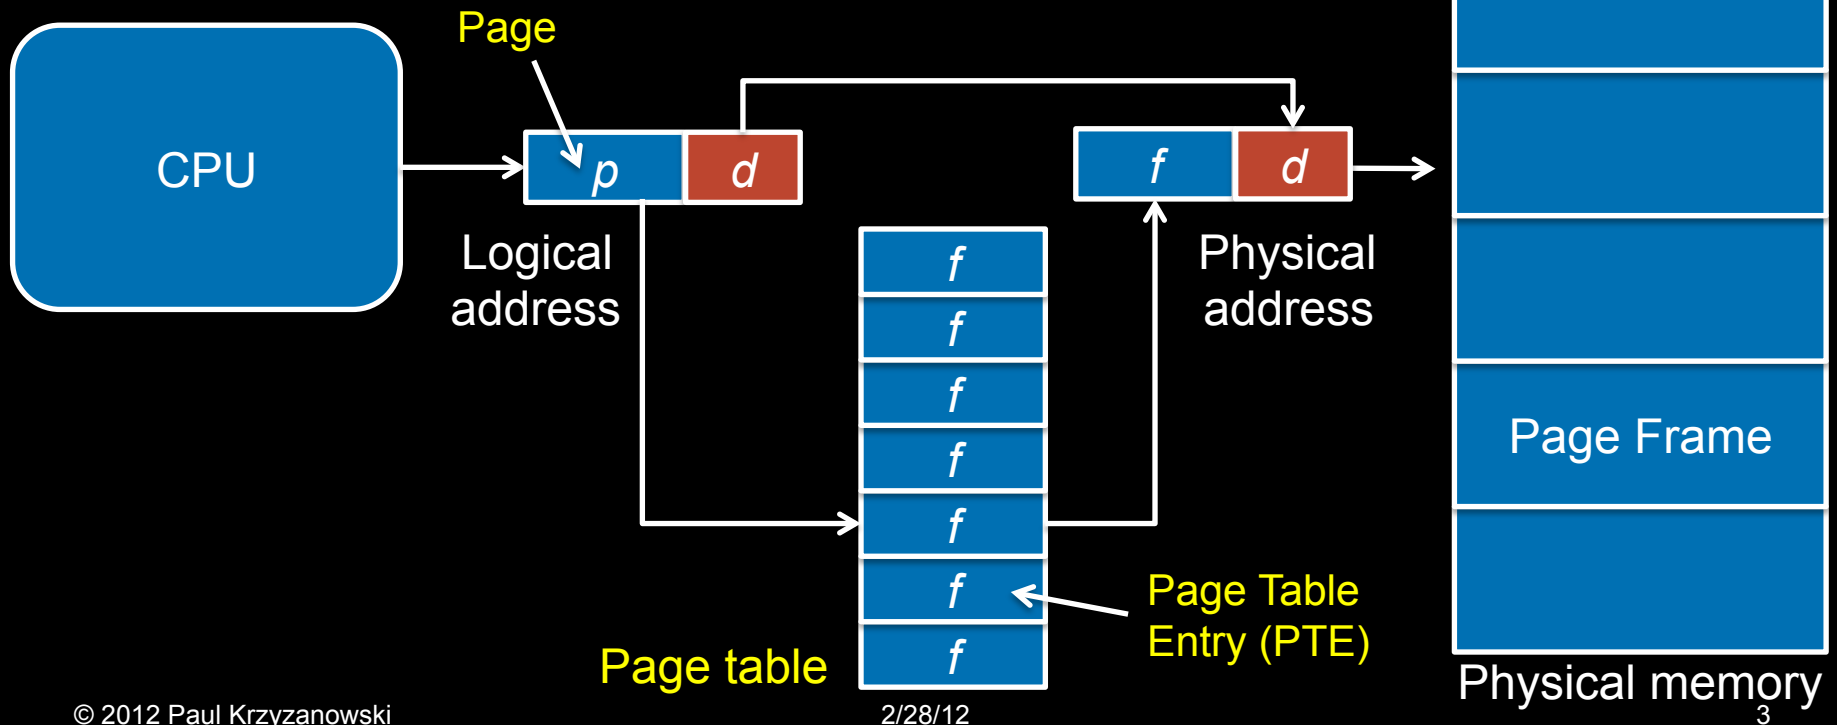
Paul Krzyzanowski
pxk@cs.rutgers.edu

Recap

Page translation



$f = \text{page_table}[p]$



Page table

- One page table per process
- Stores corresponding page frame # for a page #
- And stores page permissions:
 - Read-only
 - No-execute
 - Dirty (modified)
 - Referenced
 - Others (e.g., secure or privileged mode access)
- Page table is selected by setting a **page table base register** with the address of the table

Improving look-up performance: TLB

- Cache frequently-accessed pages
 - Translation lookaside buffer (TLB)
 - Associative memory: key (page #) and value (frame #)
- TLB is on-chip & fast ... but small (64 – 1,024 entries)
- TLB miss: result not in the TLB
 - Need to do page table lookup in memory
- Hit ratio = % of lookups that come from the TLB
- Address Space Identifier (ASID): share TLB among address spaces

Page-Based Virtual Memory Benefits

- Simplify memory management for multiprogramming
 - Allow **discontiguous allocation**
 - MMU give the illusion of contiguous allocation
- Allow a process to feel that it **has more memory** than it really has
 - Also: process can have greater address space than system memory
- Memory **Protection**
 - Each process' address space is separate from others
 - MMU allows pages to be protected:
 - Writing, execution, kernel vs. user access

Real-Time Considerations

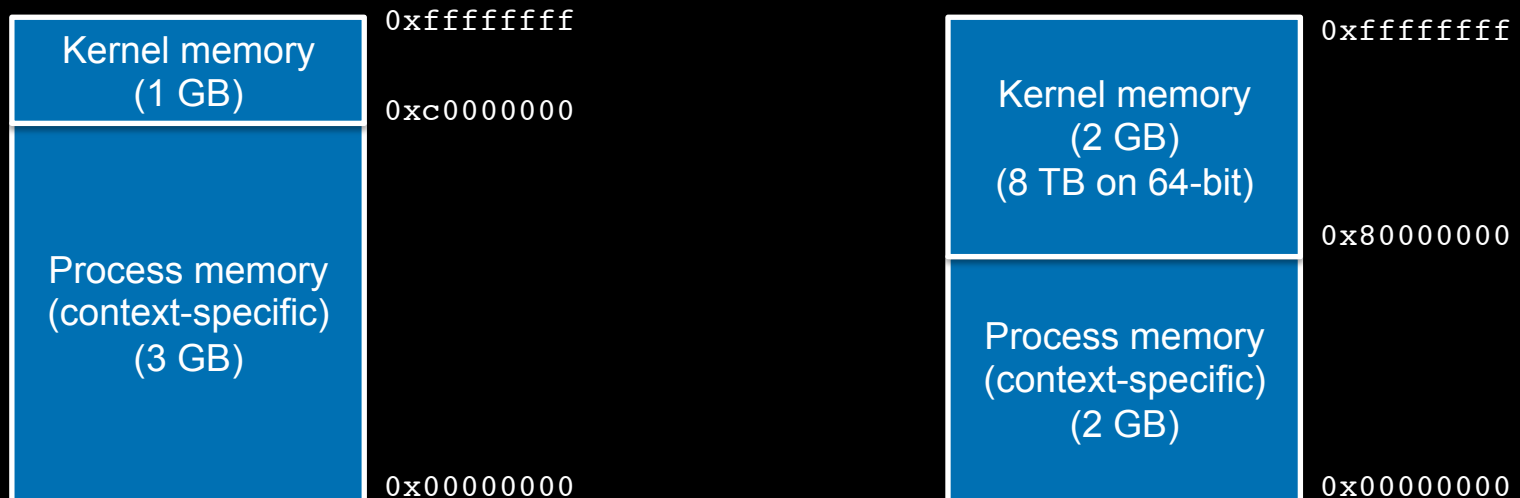
- Avoid page table lookup
 - Or run CPU without virtual addressing
- Pin high-priority real-time process memory into TLB (if possible)

Accessing memory

- Process makes *virtual* address references for all memory access
- MMU converts to physical address via a per-process page table
 - Page number → Page frame number
 - Basic info stored in a PTE (page table entry):
 - Valid? Is the page mapped?
 - Page frame number
 - Permissions (read-only, read-write, execute-only, ...)
 - Modified?
 - **Page fault** if not a valid reference
- Most CPUs support:
 - **Virtual addressing mode** and **Physical addressing mode**
 - CPU starts in physical mode ... someone has to set up page tables
 - Divide address space into user & kernel spaces

Kernel's View

- Each process sees a flat linear address space
 - Accessing regions of memory mapped to the kernel causes a page fault
- Kernel's view:
 - Address space is split into two parts
 - User part: changes with context switches
 - Kernel part: remains constant
 - Split is configurable:
 - 32-bit x86: PAGE_OFFSET: 3GB for process + 1 GB kernel



Page allocator

- With VM, processes can use non-contiguous pages
- Sometimes you need contiguous allocation
- E.g., DMA logic ignores paging
 - If we rely on DMA, we need contiguous pages

Page allocator

- Linux kernel support for contiguous buffers
 - free_area: keep track of lists of free pages
 - 1st element: free single pages
 - 2nd element: free blocks of 2 contiguous pages
 - 3rd element: free blocks of 4 contiguous pages
 - ...
 - 10th element: free blocks of 512 contiguous pages

Buddy System

- Try to get the best usable allocation unit
- If not available, get the next biggest one & split
- Coalesce upon free
- Example
 - We want 8 contiguous pages
 - Do we have a block of 8? Suppose no.
 - Do we have a block of 16? Suppose no.
 - Do we have a block of 32? Suppose yes.
 - Split the 32 block into two blocks of 16. Back up.
 - Do we have a block of 16? Yes!
 - Split one of the 16 blocks into two blocks of eight. Back up.
 - Do we have a block of 8? Yes!

Buddy System: Coalescence

- When a block is freed, see if we can merge buddies
- Two blocks are buddies if:
 - They are the same size, b
 - They are contiguous
 - The address of the first page of the lower # block is a multiple of $2b \times \text{page_size}$
- If two blocks are buddies, they are merged
- Repeat the process.

Buddy System Example



512 blocks

256 blocks

128 blocks

64 blocks

We want a 64-block allocation.


None available.

Any 128-block chunks to split? No.

Any 256-block chunks to split? No.

Any 512-block chunks to split? Yes.

Buddy System Example



512 blocks

256 blocks

128 blocks


64 blocks

Split a 512-block chunk into two 256-block chunks.
Try again.

We want a 64-block allocation.
None available.

Any 128-block chunks to split? No.
Any 256-block chunks to split? Yes.

Buddy System Example



512 blocks

256 blocks

128 blocks


64 blocks

Split a 256-block chunk into two 128-block chunks.
Try again.

We want a 64-block allocation.
None available.

Any 128-block chunks to split? Yes.

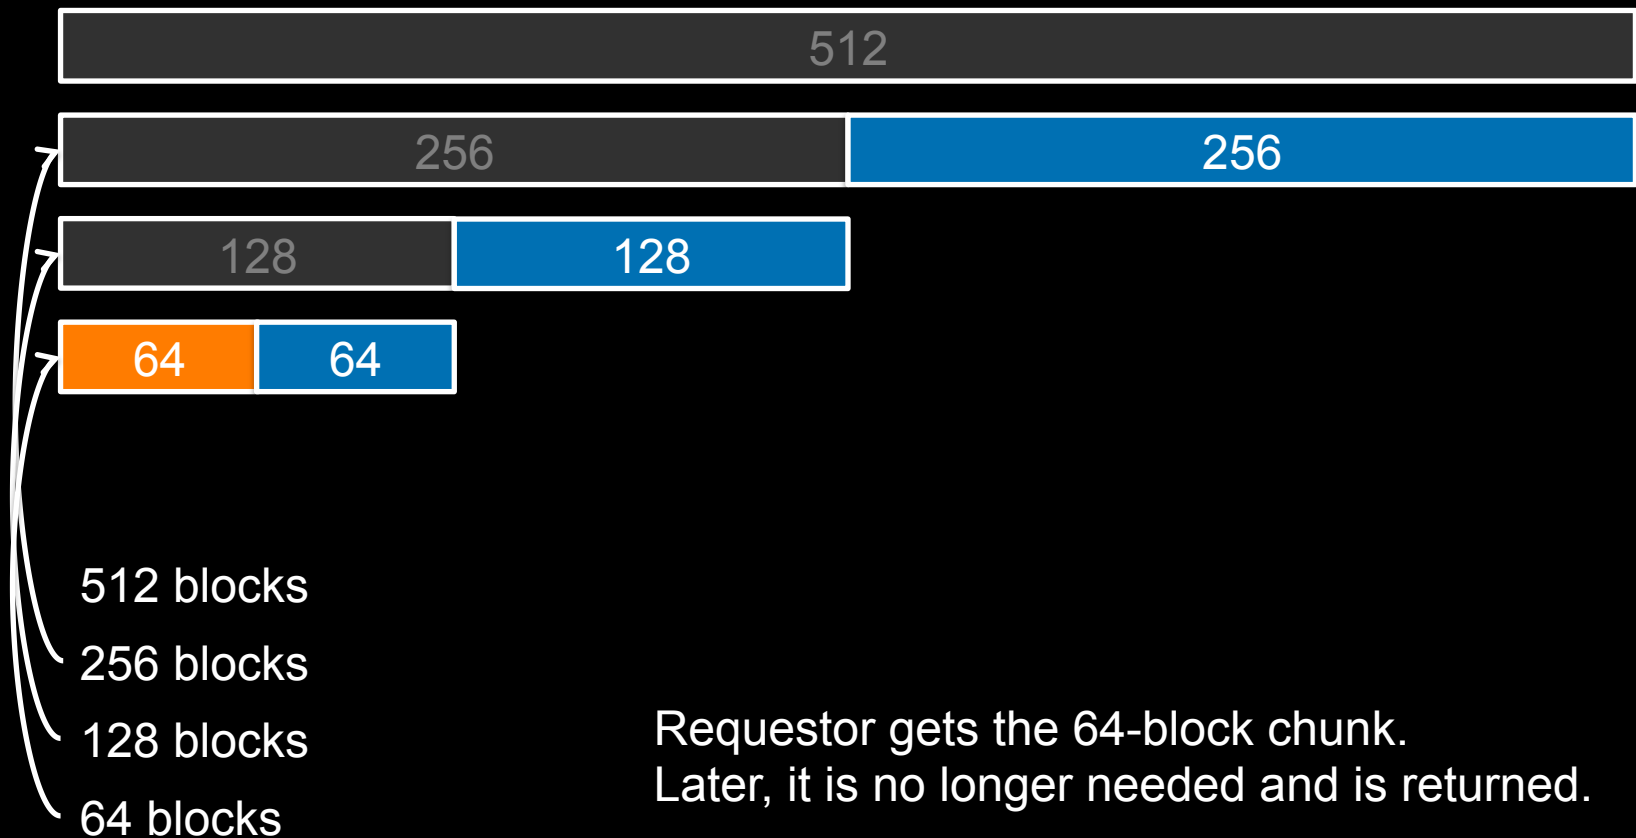
Buddy System Example




Split a 128-block chunk into two 64-block chunks.
Try again.

We want a 64-block allocation.
Got it!

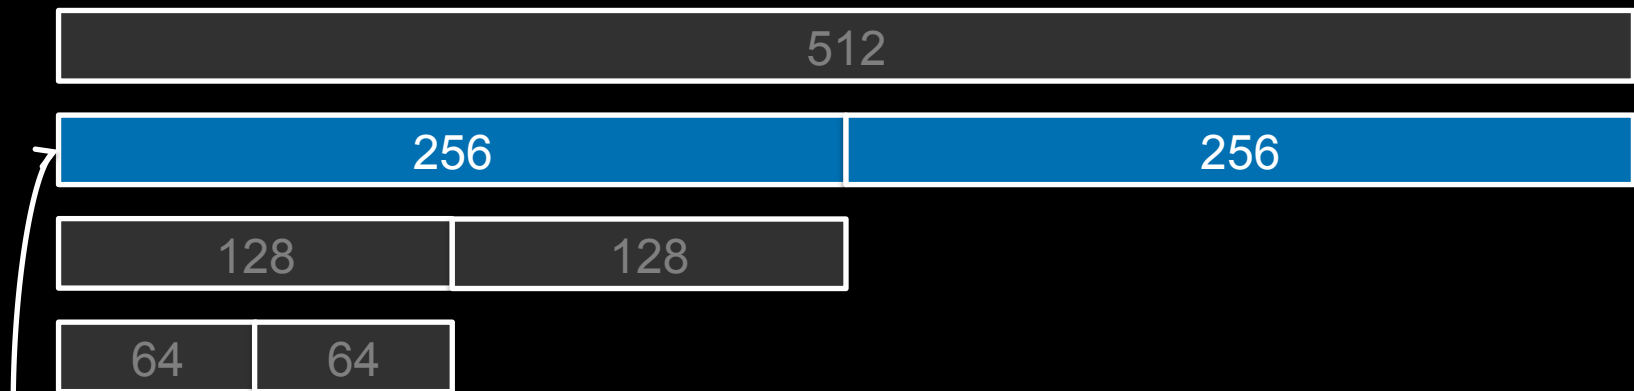
Buddy System Example



Buddy System Example



Buddy System Example



512 blocks

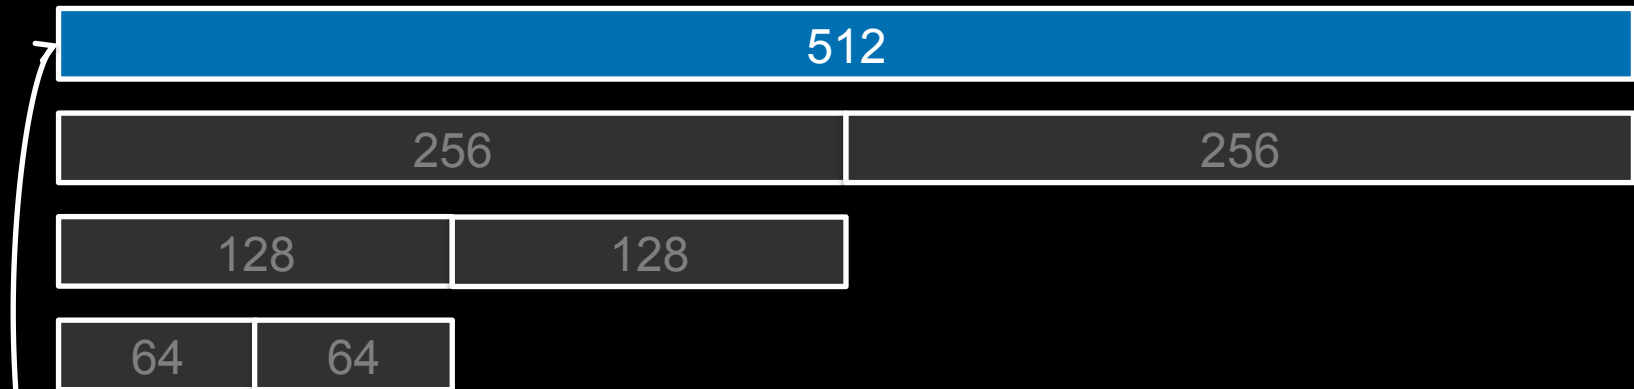
256 blocks

128 blocks

64 blocks

Combine the two 128-block chunks into a 256-block chunk

Buddy System Example



512 blocks

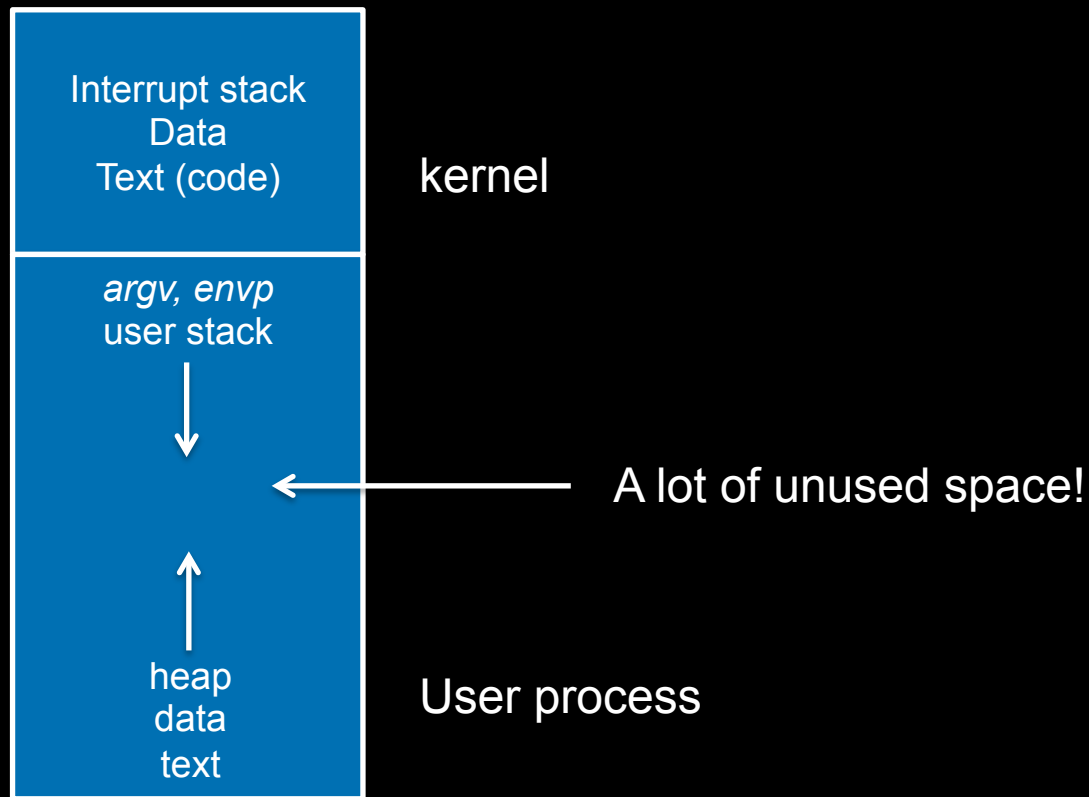
256 blocks

128 blocks

64 blocks

Combine the two 256-block chunks into a 512-block chunk


Sample memory map per process




Multilevel (Hierarchical) page tables

- Most processes use only a small part of their address space
- Keeping an entire page table is wasteful
- E.g., 32-bit system with 4KB pages: 20-bit page table
 - 1,048,576 entries in a page table

Multilevel page table

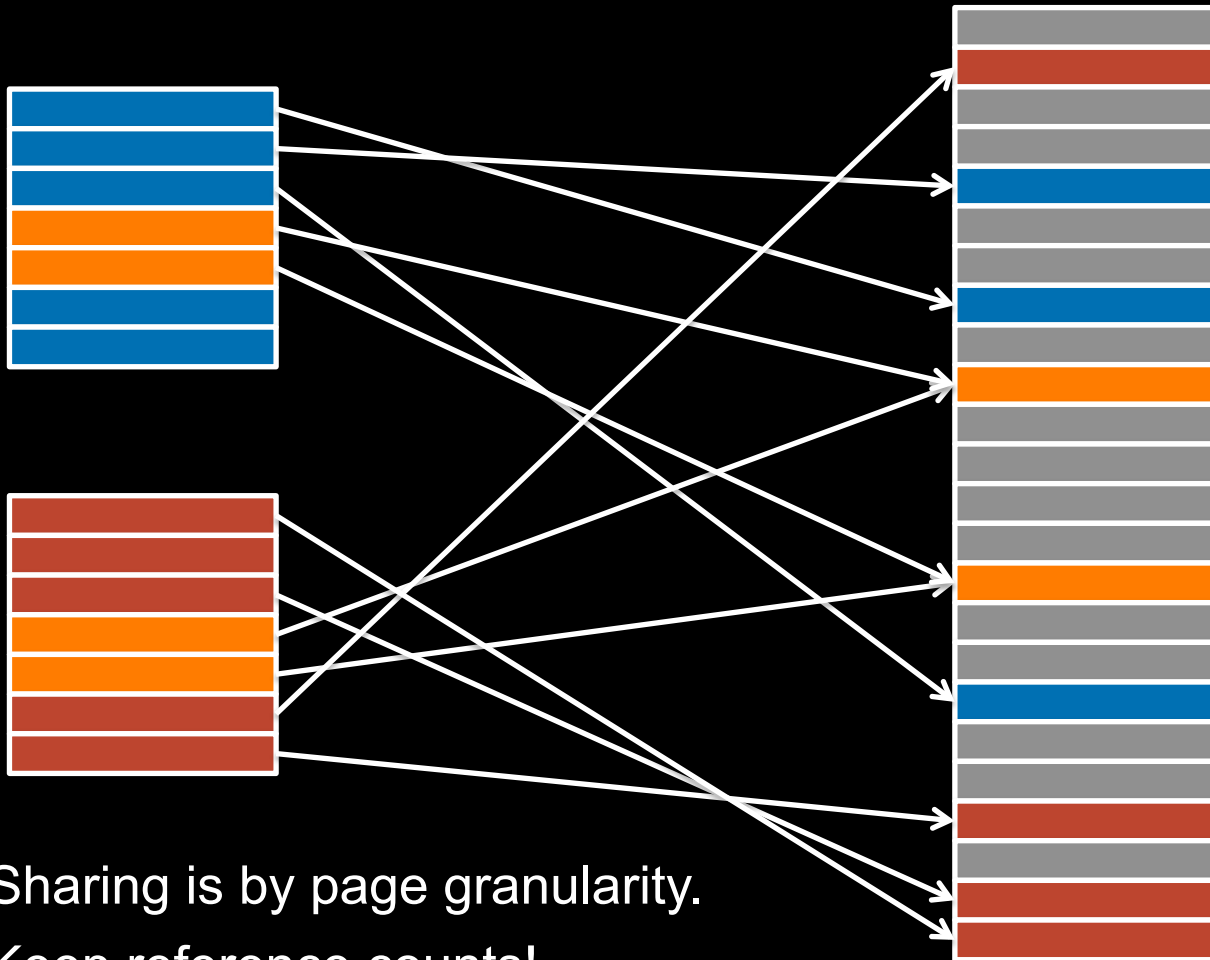


Virtual memory makes memory sharing easy



Sharing is by page granularity.

Virtual memory makes memory sharing easy



Sharing is by page granularity.
Keep reference counts!

Copy on write

- Share until a page gets modified
- Example: `fork()`
 - Set all pages to read-only
 - Trap on write
 - If legitimate write
 - Allocate a new page and copy contents

MMU Example: ARM

ARMv7-A architecture

- Cortex-A8
 - iPhone 3GS, iPod Touch 3G, Apple A4 processor in iPhone 4 & iPad, Droid X, Droid 2, etc.)
- Cortex-A9
 - Multicore support
 - TI OMAP 44xx series, Apple A5 processor in iPad 2

Pages

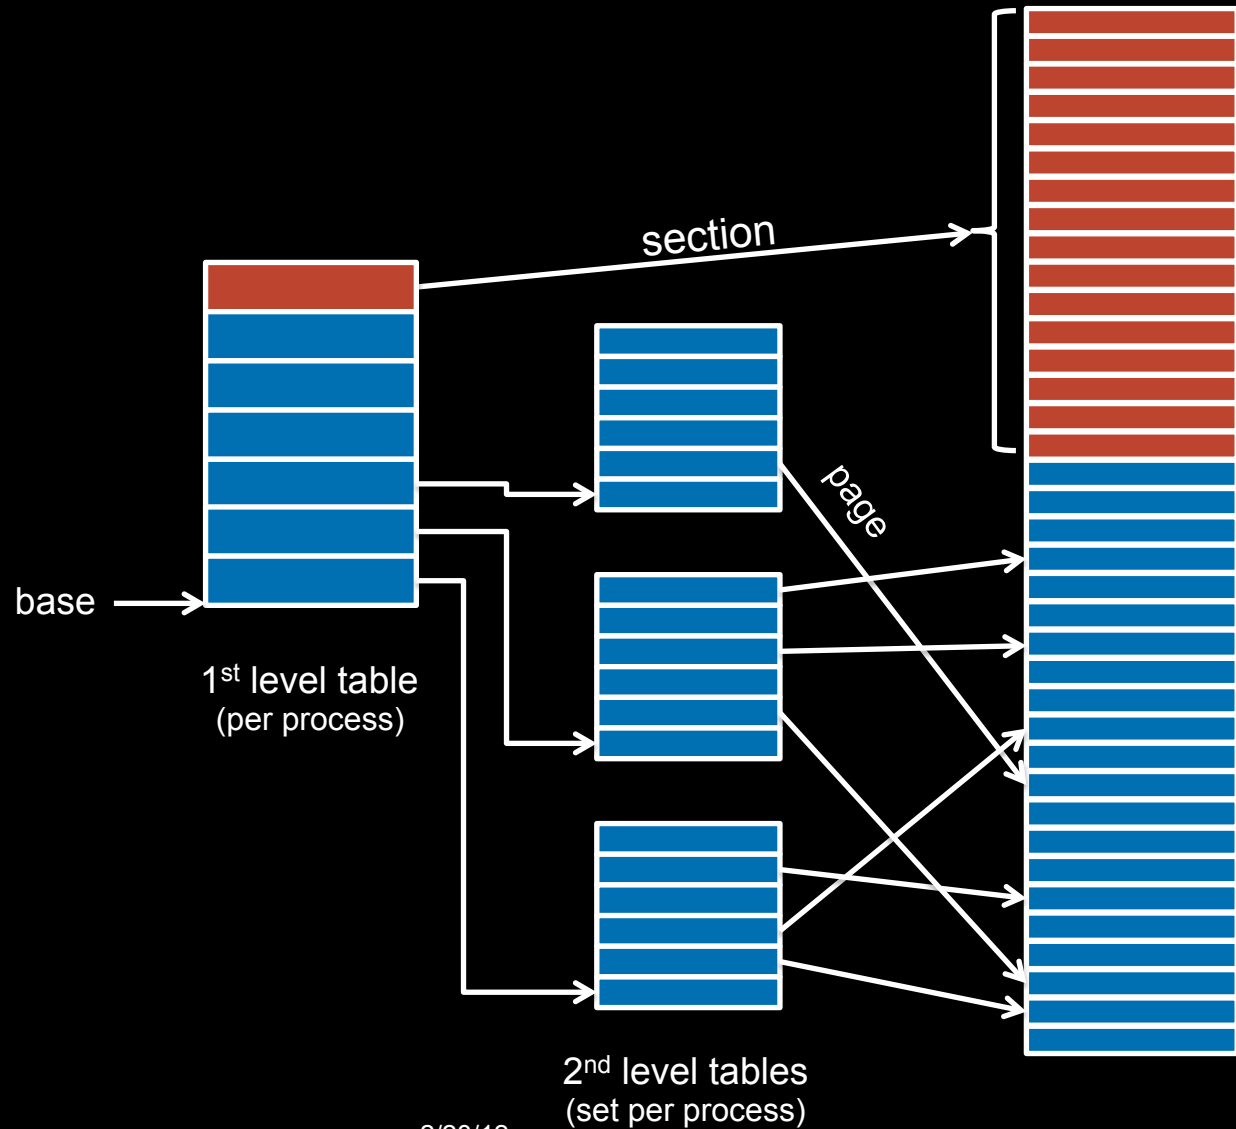
Four page (block) sizes:

- Supersections: 16MB memory blocks
- Sections: 1MB memory blocks
- Large pages: 64KB memory blocks
- Small pages: 4KB memory blocks

Two levels of tables

- **First level table**
 - Base address, descriptors, and translation properties for sections and supersections (1 MB & 16 MB blocks)
 - Translation properties and pointers to a **second level table** for large and small pages (4 KB and 64 KB pages)
- **Second level tables** (*aka* page tables)
 - Each contains base address and translation properties for small and large pages
- Benefit: a large region of memory can be mapped using a single entry in the TLB (e.g., OS)

ARM Page Tables



TLB

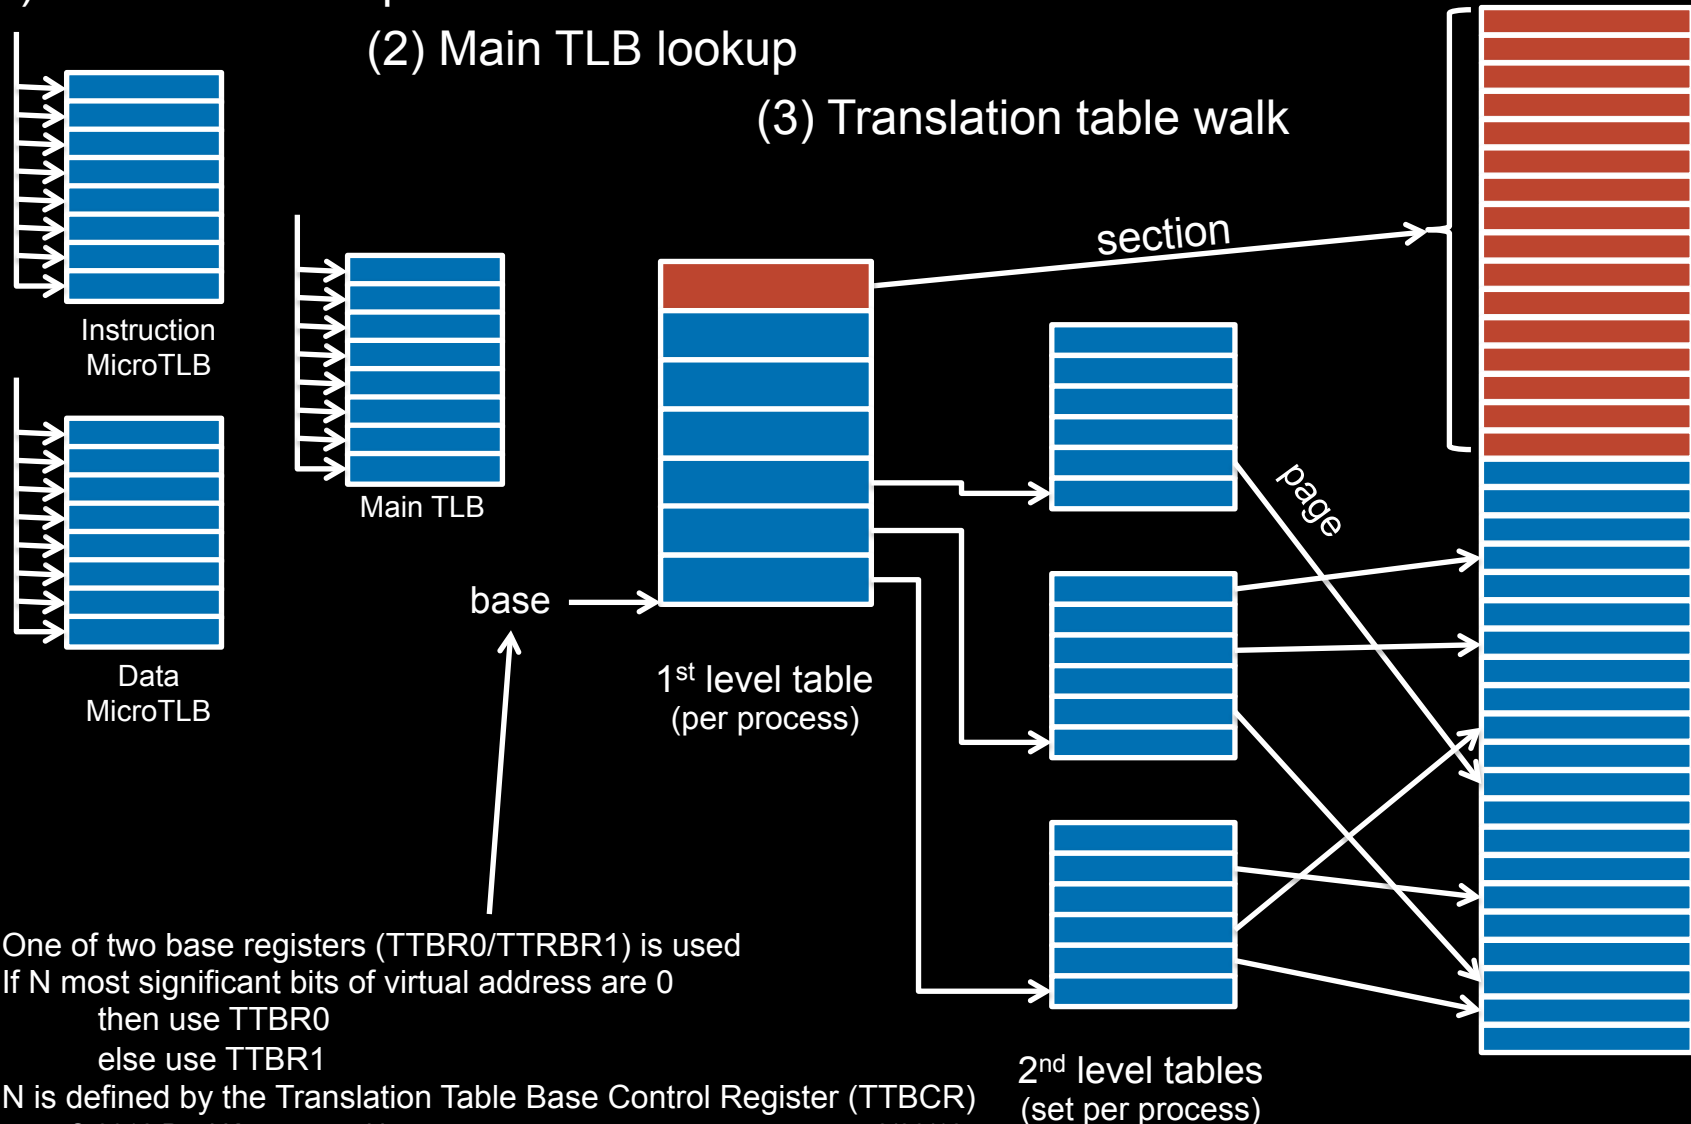
- 1st level: **MicroTLB** – one each for instruction & data sides
 - 32 entries (10 entries in older v6 architectures)
 - Address Space Identifier (**ASID**) [8 bits] and Non-Secure Table Identifier (NSTID) [1 bit]; entries can be global
 - Fully associative; one-cycle lookup
 - Lookup checks protection attributes: may signal Data Abort
 - Replacement either Round-Robin (default) or Random
- 2nd level: **Main TLB** – catches cache misses from microTLBs
 - 8 fully associative entries (may be locked) + 64 low associative entries
 - variable number of cycles for lookup
 - lockdown region of 8 entries (important for real-time)
 - Entries are globally mapped or associated ASID and NSTID

ARM Page Tables

(1) MicroTLB lookup

(2) Main TLB lookup


(3) Translation table walk



One of two base registers (TTBR0/TTBR1) is used
If N most significant bits of virtual address are 0
then use TTBR0
else use TTBR1


N is defined by the Translation Table Base Control Register (TTBCR)

Translation flow for a section (1 MB)



Physical section = Read [Translation table base + table index]

Physical address = physical section : section offset



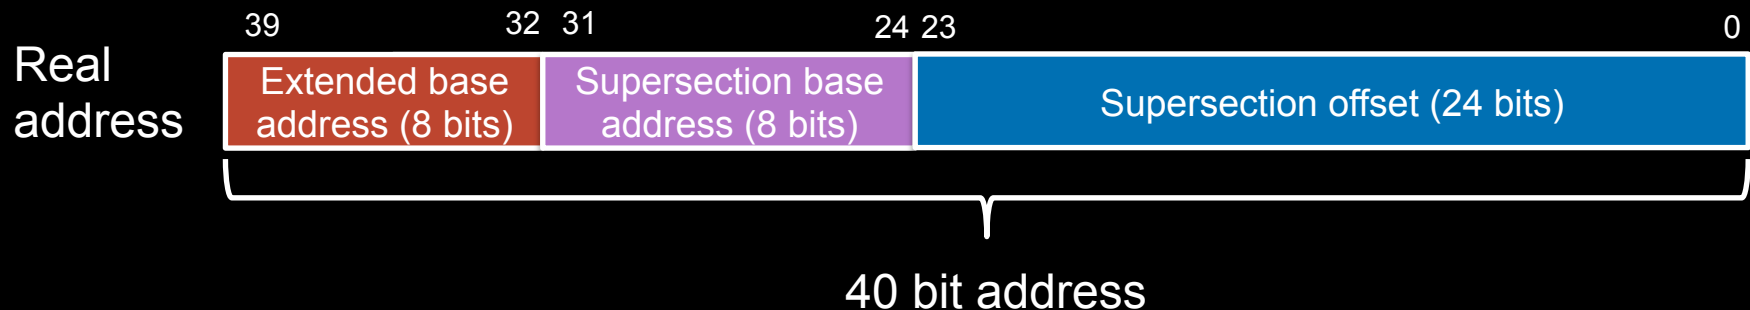
Translation flow for a supersection (16 MB)




Supersection base address, Extended base address =

Read [Translation table base + table index]

Real address = Extended base address : physical section : section offset



Translation flow for a small page (4KB)



Page table address = Read [Translation table base + first-level index]

Physical page = read[page table address + second-level index]

Real address = physical page : page offset



Translation flow for a large page (64KB)



Page table address = Read [Translation table base + first-level index]

Physical page = read[page table address + second-level index]

Physical address = physical page : page offset



Memory Protection & Control

- Domains
 - Clients execute & access data within a **domain**. Each access is checked against access **permissions** for each memory block
- Memory region attributes
 - Execute never
 - Read-only, read/write, no access
 - Privileged read-only, privileged & user read-only
 - Non-secure (is this secure memory or not?)
 - Sharable (is this memory shared with other processors)
 - Strongly ordered (memory accesses must occur in program order)
 - Device/shared, device/non-shared
 - Normal/shared, normal/non-shared
- Signal *Memory Abort* if permission is not valid for access

MMU Example: x86-64

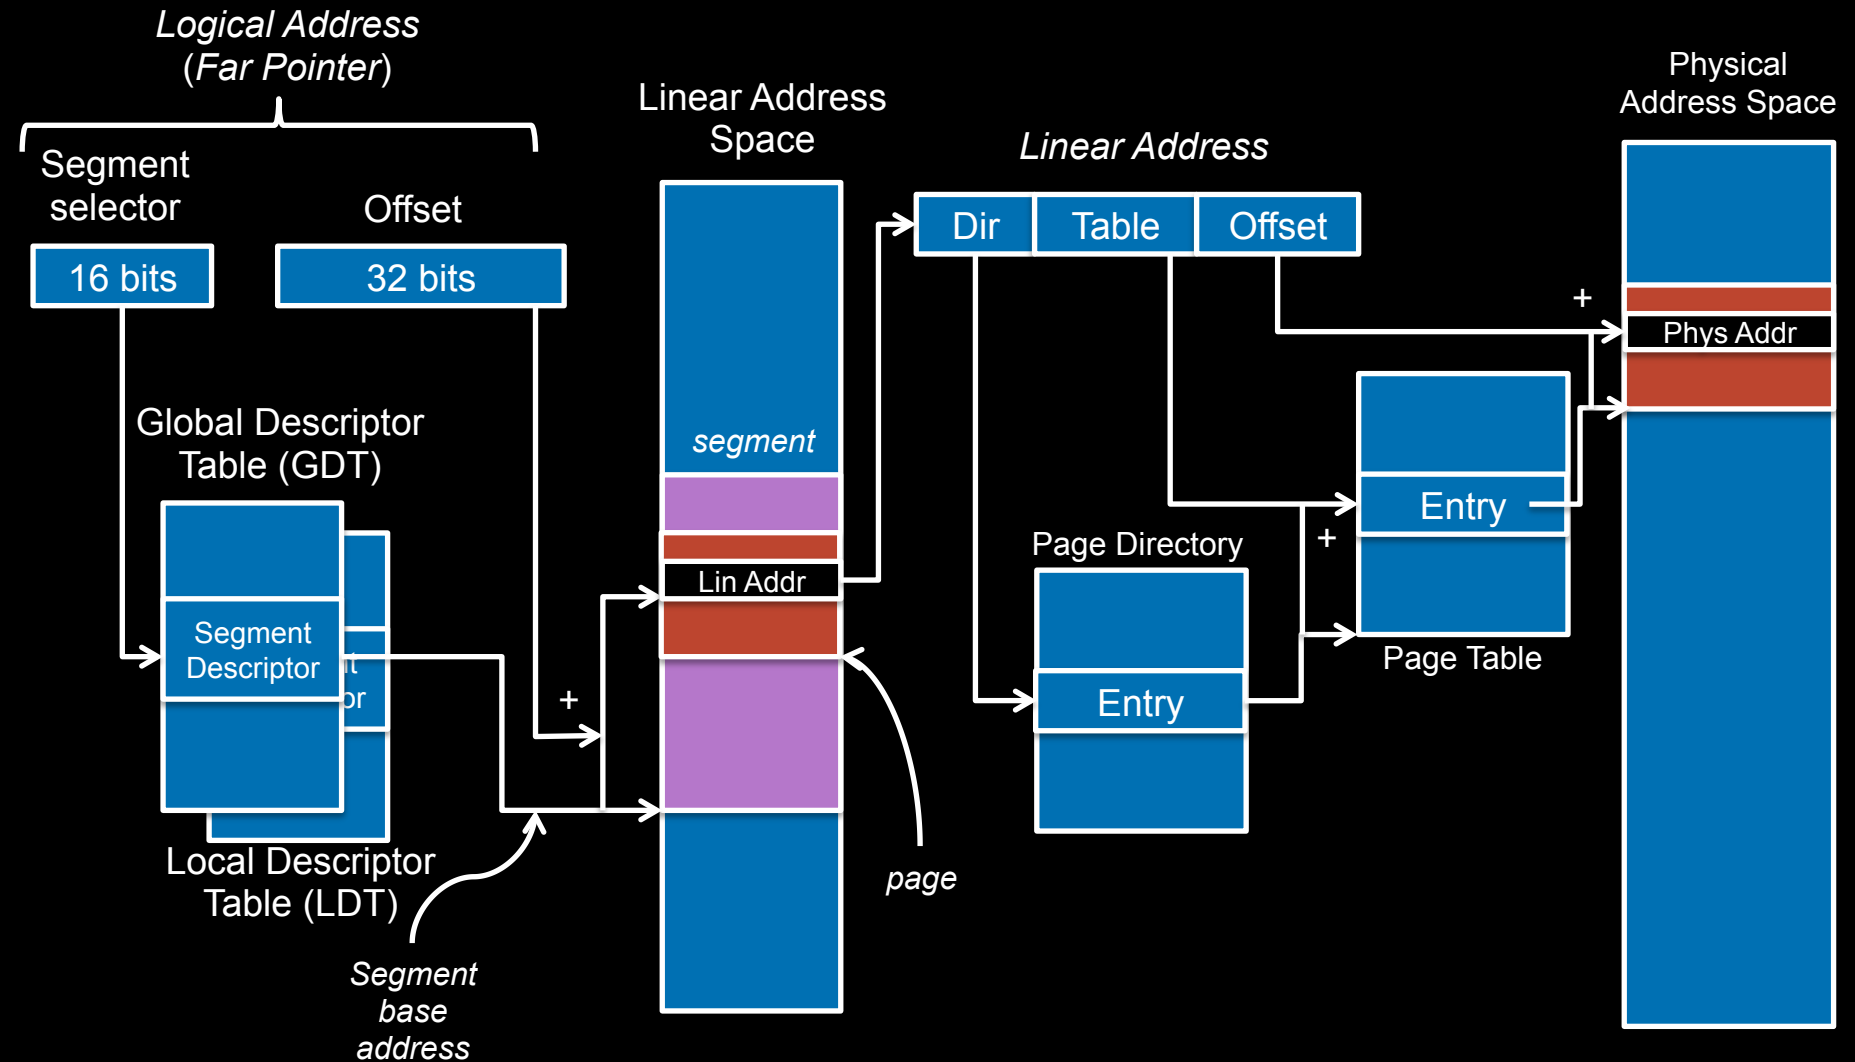
IA-32 Memory Models

- Flat memory model
 - Linear address space
 - Single, contiguous address space
- Segmented memory model
 - Memory appears as a group of independent address spaces: segments (code, data, stack, etc.)
 - Logical address = {segment selector, offset}
 - 16,383 segments; each segment can be up to 2^{32} bytes
- Real mode
 - 8086 model
 - Segments up to 64KB in size
 - maximum address space: 2^{20} bytes

Segments

- Each segment may be up to 4 GB
- Up to 16 K segments per process
- Two partitions per process
 - Local: *private to the process*
 - Up to 8 K segments
 - Info stored in a Local Descriptor Table (LDT)
 - Global: *shared among all processes*
 - Up to 8 K segments
 - Info stored in a Global Descriptor Table (GDT)
- Logical address is (*segment selector, offset*)
 - Segment selector = 16 bits:
 - 13 bits segment number + 1 bit LDT/GDT ID + 2 bits protection

IA-32 Segmentation & Paging



Segment protection

- S flag in segment descriptor identifies *code* or *data* segment
- Accessed (referenced)
 - has the segment been accessed since the last time the OS cleared the bit?
- Dirty
 - Has the page been modified?
- Data
 - Write-enable
 - Read-only or read/write?
 - Expansion direction
 - Expand down (e.g., for stack): dynamically changing the segment limit causes space to be added to the bottom of the stack
- Code
 - Execute only, execute/read (e.g., constants in code segment)
 - Conforming:
 - Execution can continue even if privilege level is elevated

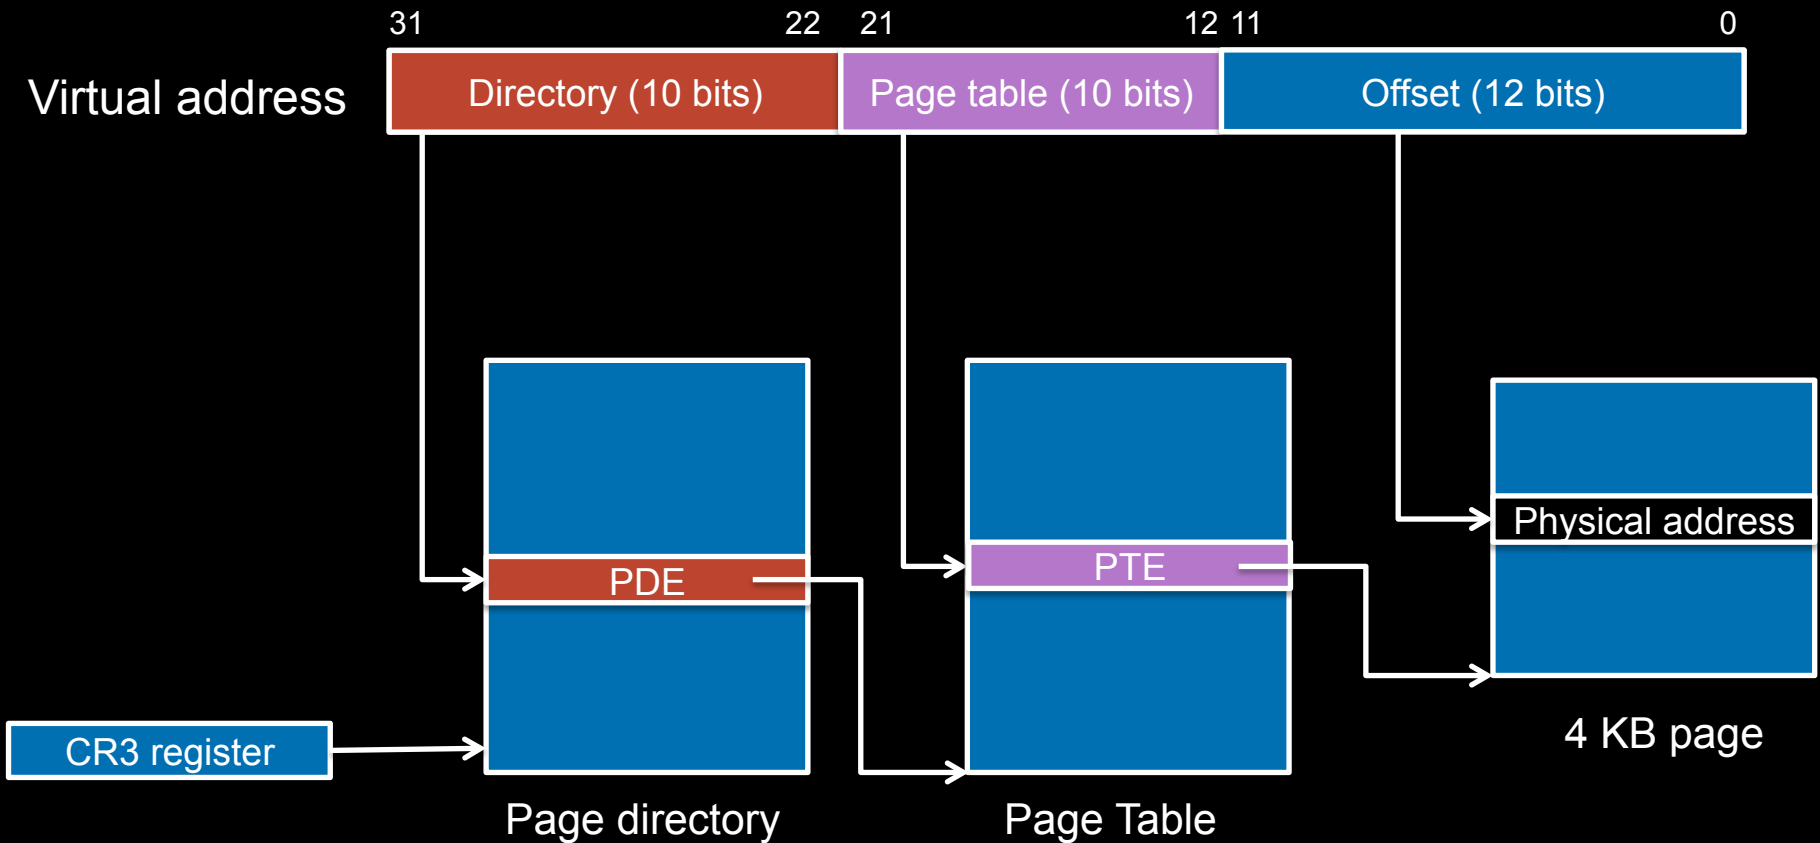
IA-32 Paging

- 32-bit registers, 36-bit address space (64 GB)
 - Physical Address Extension (PAE)
 - Bit 5 of control register CR4
 - 52 bit physical address support (4 PB of memory)
 - Only a 4 GB address space may be accessed at one time
 - Page Size Extensions (PSE-36)
 - 36-bit page size extension (64 GB of memory)
 - Supports up to 4 MB pages


Intel 64-bit mode

- Segments supported only in IA-32 emulation mode
 - Mostly disabled for 64-bit mode
 - 64-bit base addresses where used
- Three paging modes
 - 32-bit paging
 - 32-bit virtual address; 32-40 bit physical
 - 4 KB or 4 MB pages
 - PAE
 - 32-bit virtual addresses; up to 52-bit physical address
 - 4 KB or 2 MB pages
 - IA-32e paging
 - 48-bit virtual addresses; up to 52-bit physical address
 - 4 KB, 2 MB, or 1 GB pages

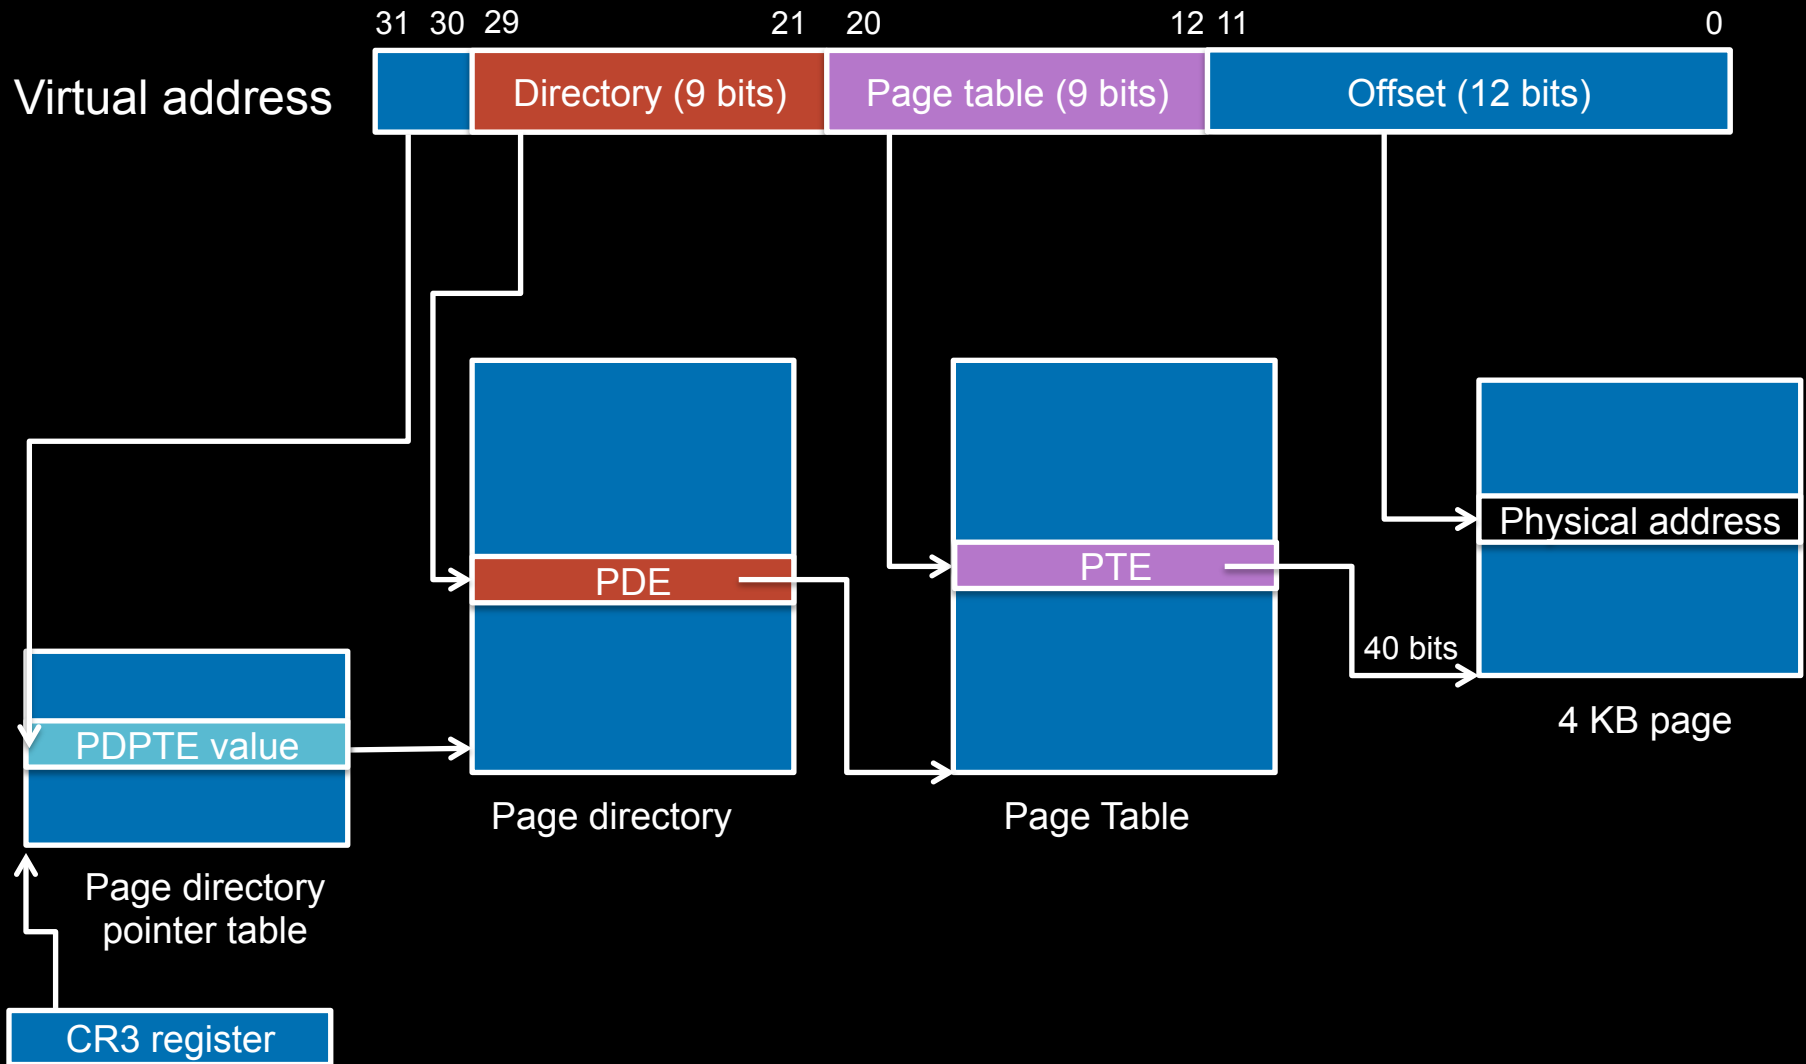
32-bit paging with 4 KB pages



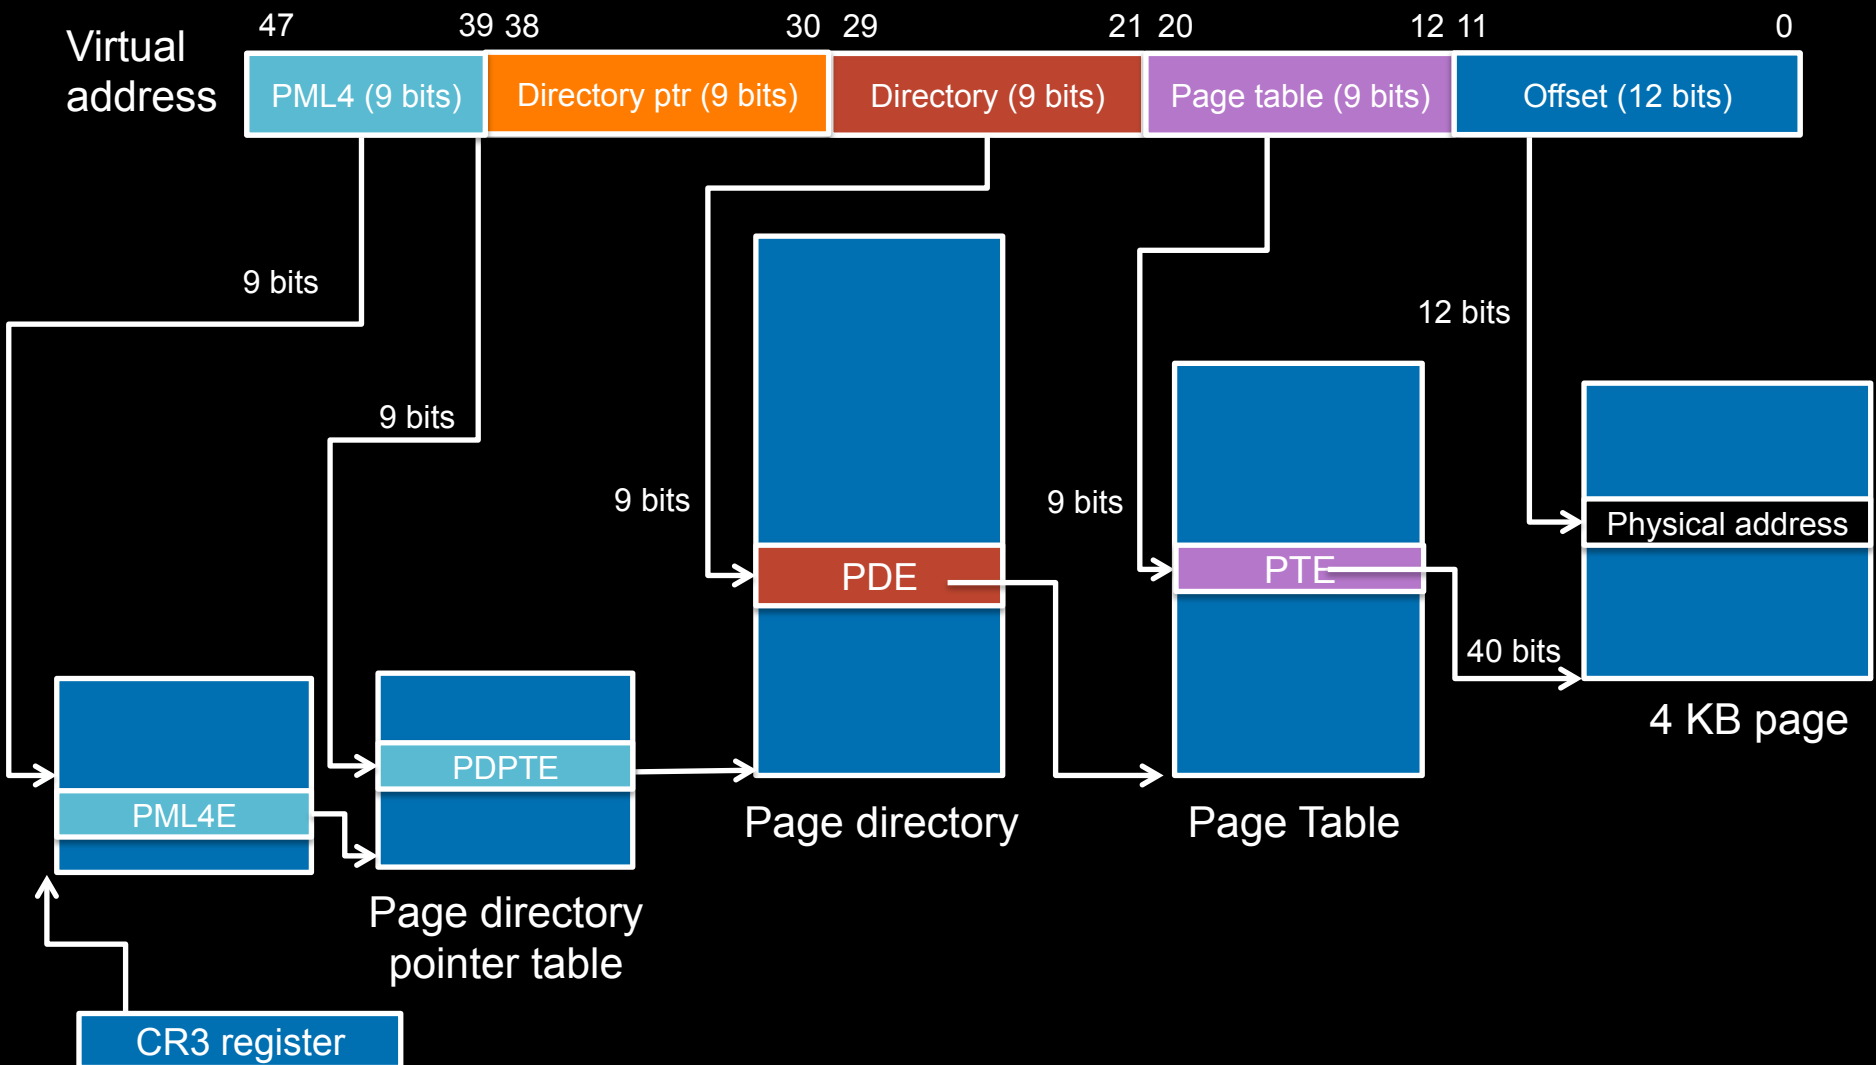
32-bit paging with 4 MB pages



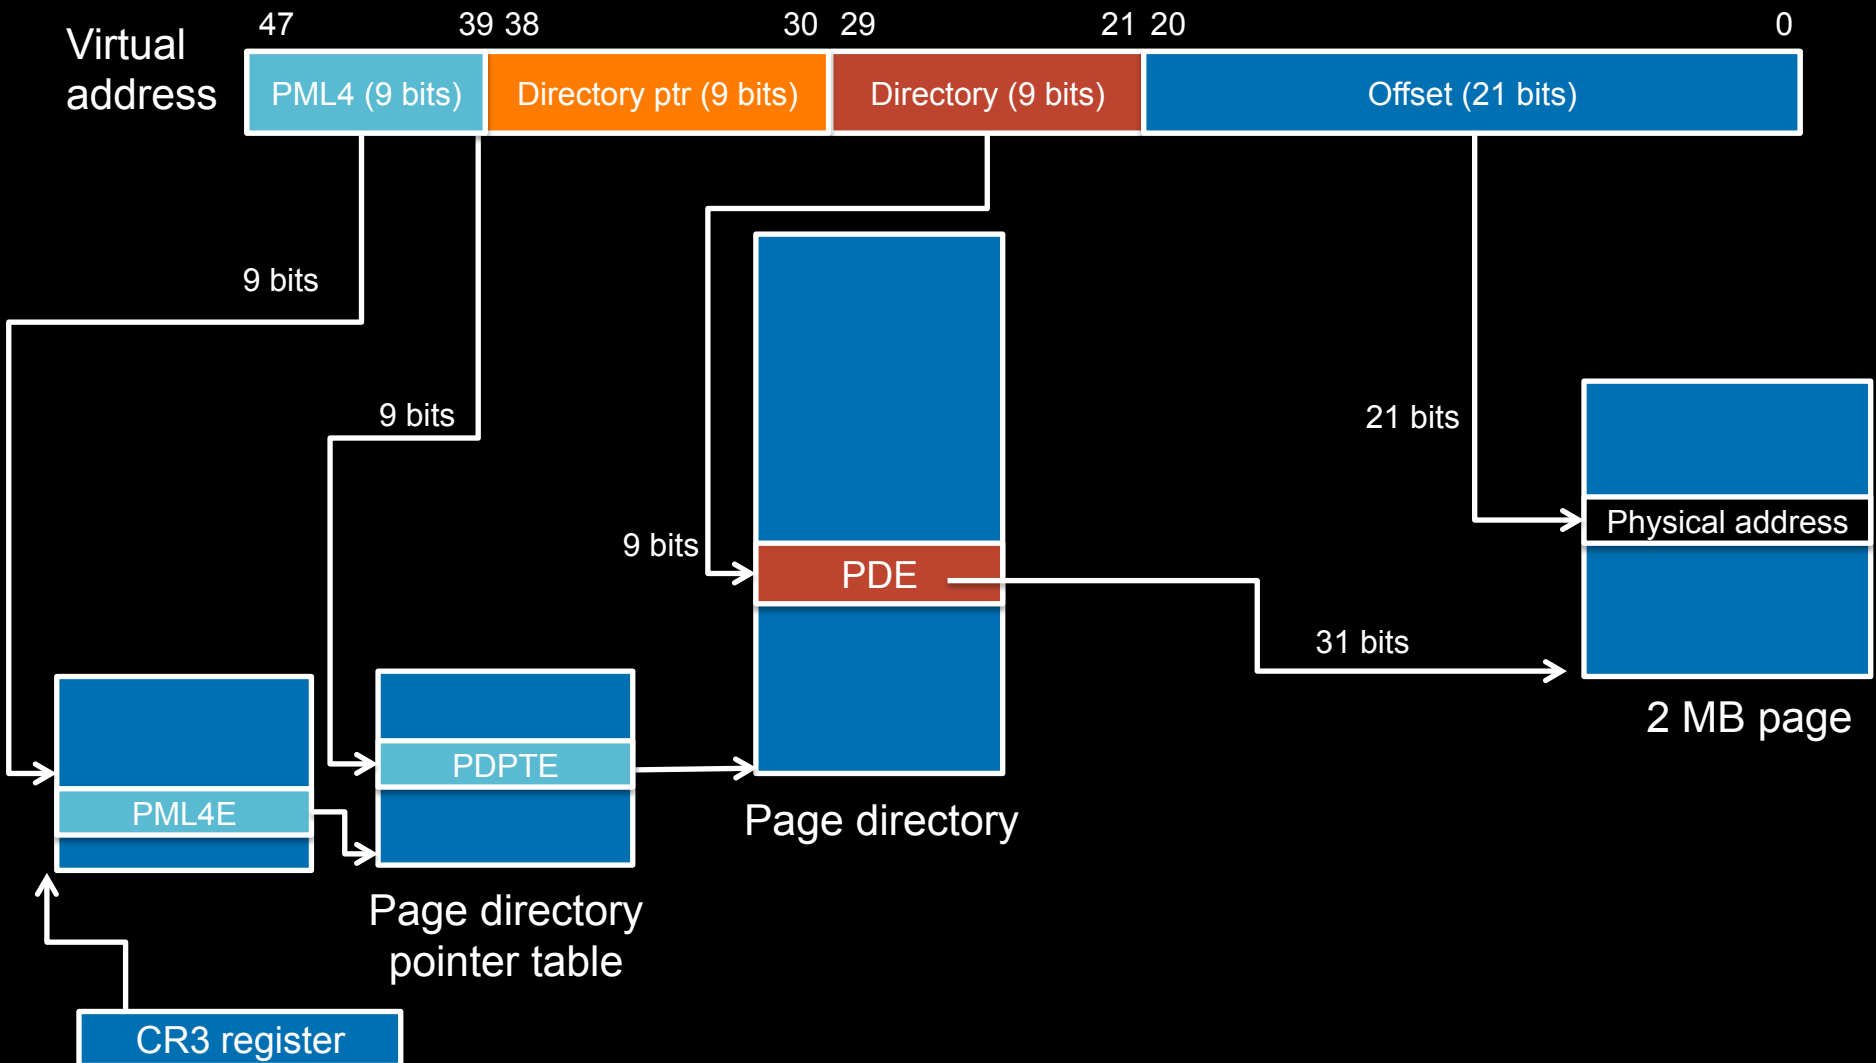
32-bit paging with 4 KB pages & PAE paging



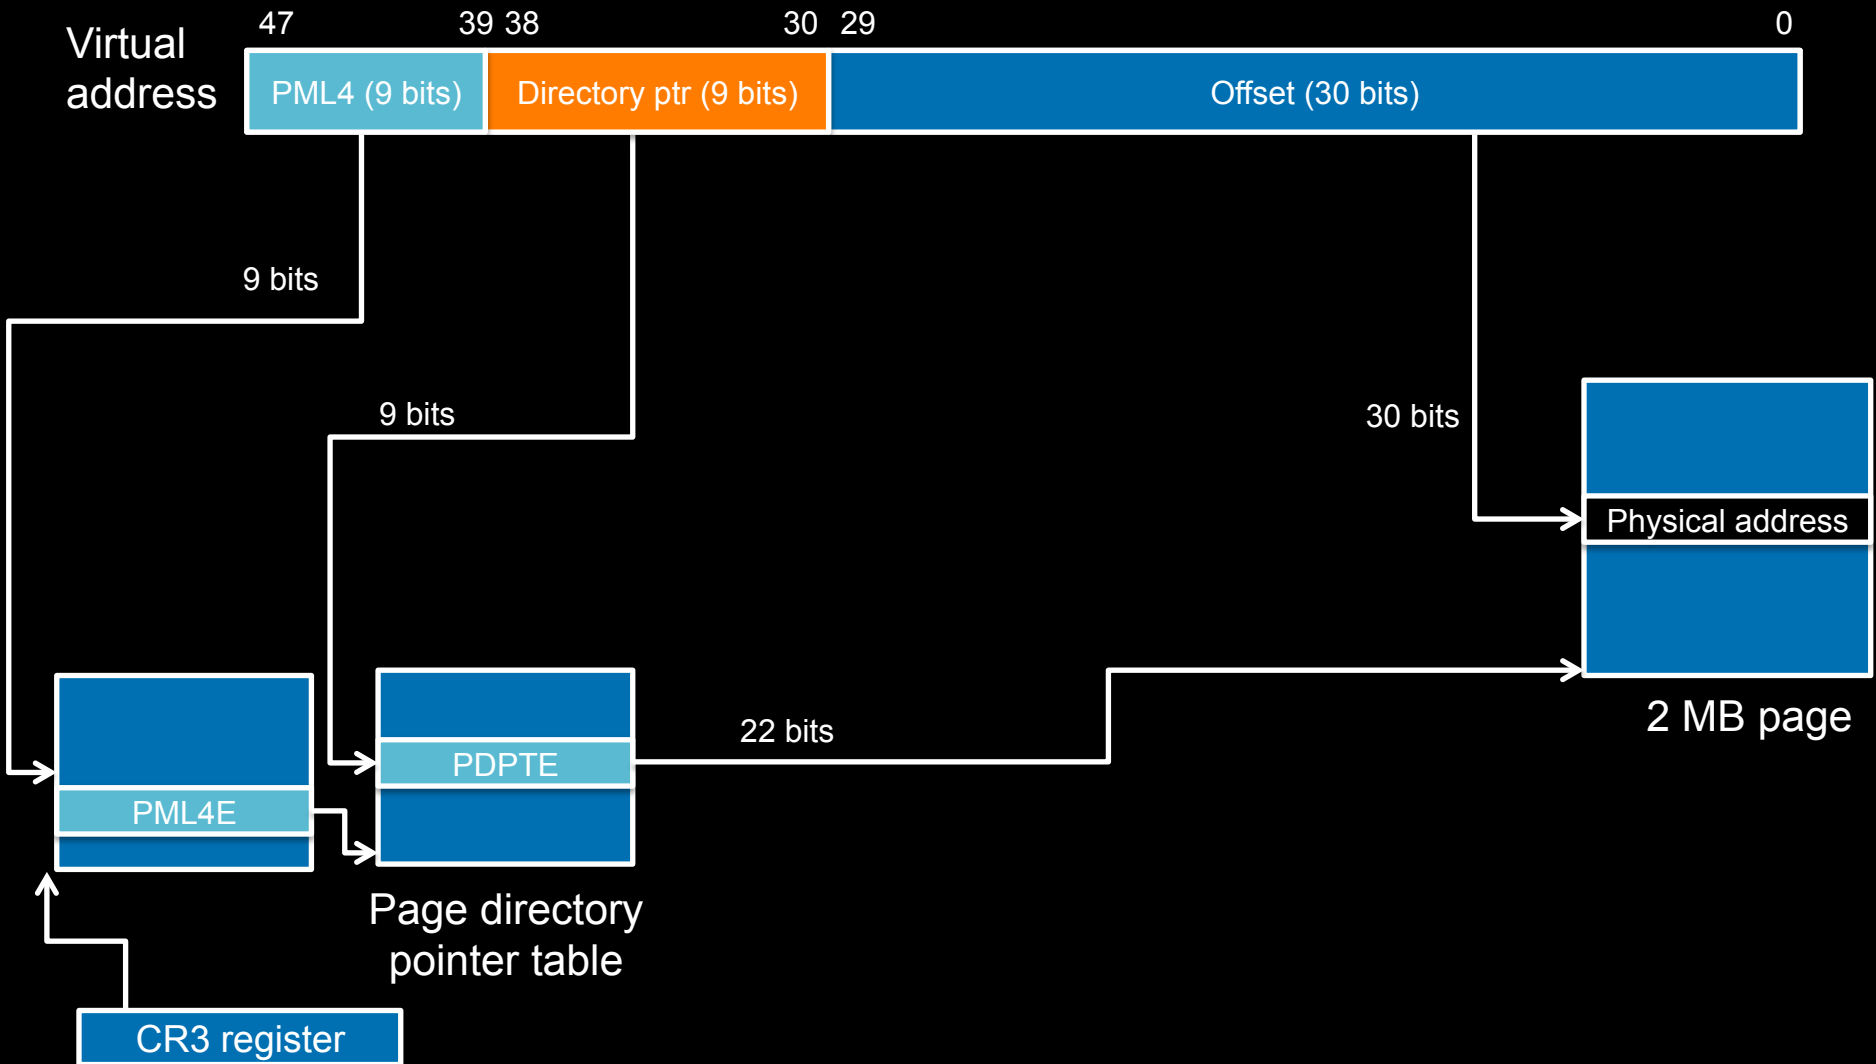
IA-32e paging with 4 KB pages



IA-32e paging with 2 MB pages



IA-32e paging with 1 GB pages



Example: TLBs on the Core i7

- 4 KB pages
 - Instruction TLB: 128 entries per core
 - Data TLB: 64 entries
 - Core 2 Duo: 16 entries TLB0; 256 entries TLB1
 - Atom: 64-entry TLB, 16-entry PDE
- Second-level unified TLB
 - 512 entries

Managing Page Tables

- Linux: architecture independent (mostly)
 - Avoids segmentation (only Intel supports it)
- Abstract structures to model 4-level page tables
 - Actual page tables are stored in a machine-specific manner

Recap

- Fragmentation is a non-issue
- Page table
- Page table entry (PTE)
- Multi-level page tables
- Inverted page table
- Segmentation
- Segmentation + Paging
- Memory protection
 - Isolation of address spaces
 - Access control defined in PTE

Demand Paging

Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it

Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it

We don't need to do this!

Demand Paging

- Load pages into memory only as needed
 - On first access
 - Pages that are never used never get loaded
- Use valid/invalid bit in page table entry
 - Valid: the page is in memory (“valid” mapping)
 - Invalid: out of bounds access **or** page is not in memory
 - Have to check the process’ memory map in the PCB to find out
- Invalid memory access generates a **page fault**

Demand Paging: At Process Start

- Open executable file
- Set up memory map (stack & text/data/bss)
 - But don't load anything!
- Load first page & allocate initial stack page
- Run it!

Memory Mapping

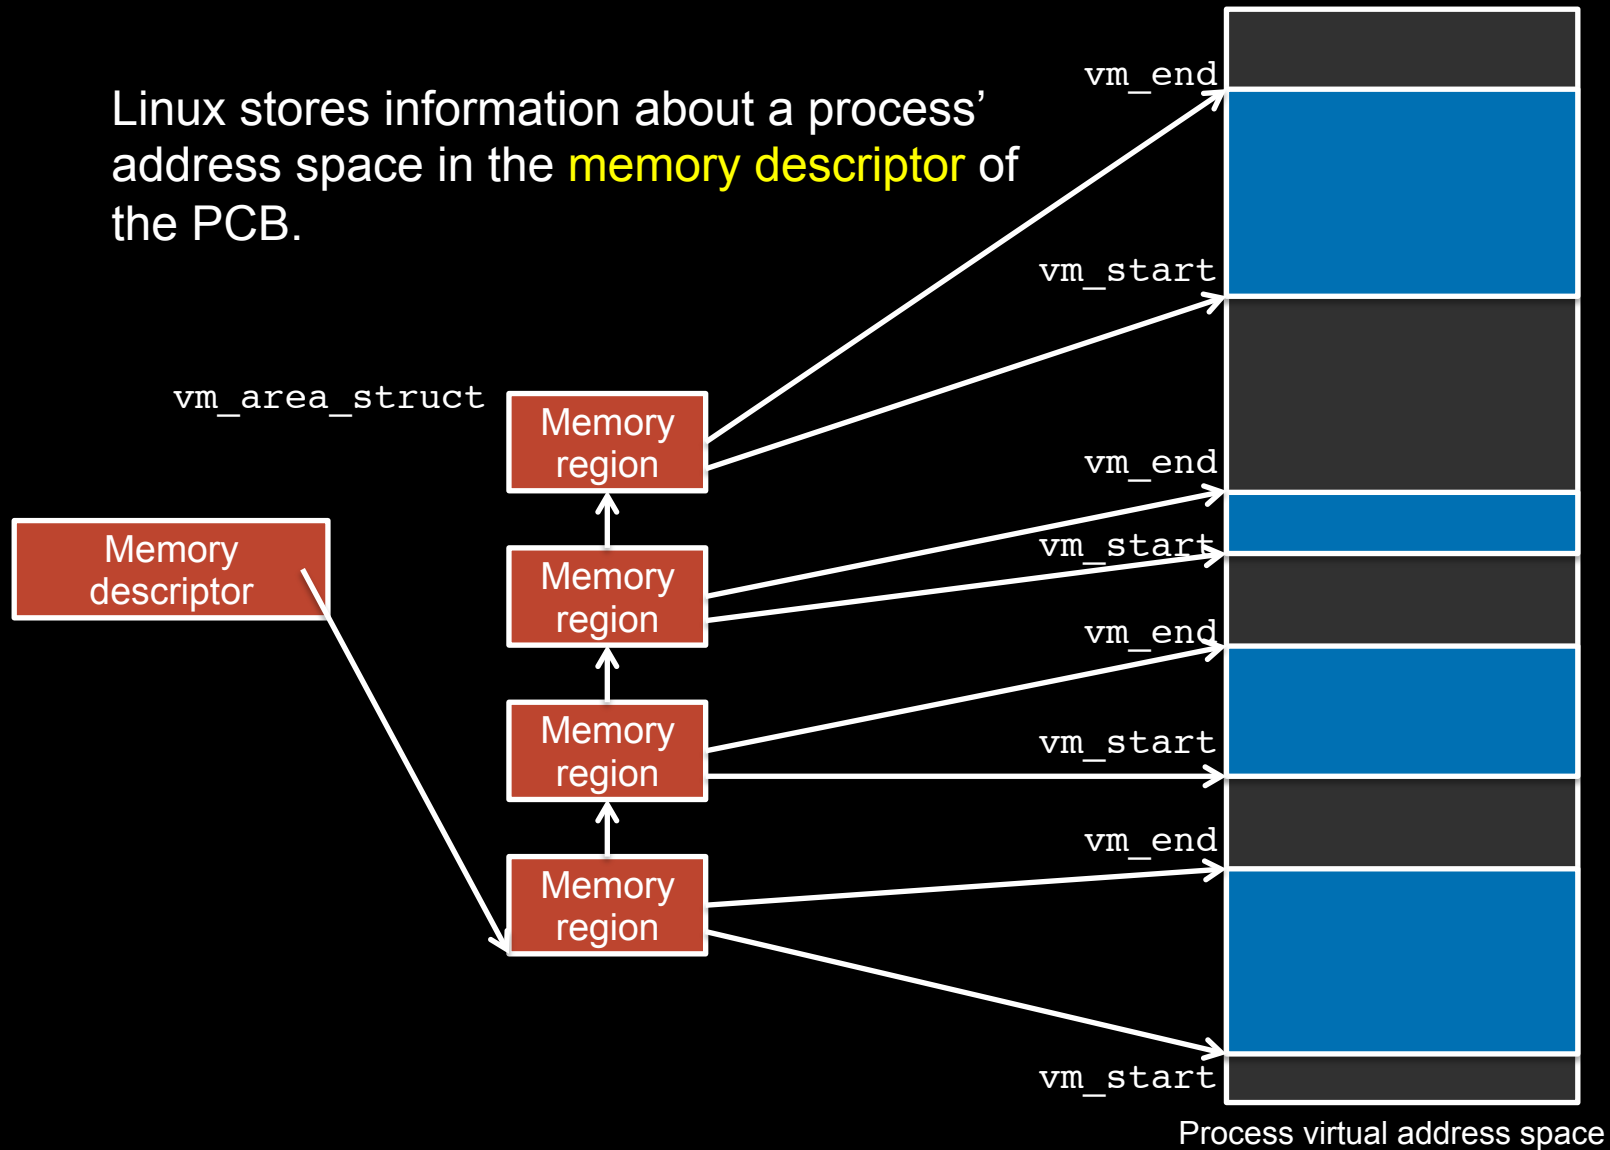
- Executable files & libraries must be brought into a process' virtual address space
 - File is *mapped* into the process' memory
 - As pages are referenced, page frames are allocated & pages are loaded into them
- If we ever run out of memory, we may need to save some modified pages into a *swap file* and load those in later on demand.
- `vm_area_struct`
 - Defines regions of virtual memory
 - Used in setting page table entries
 - Start of VM region, end of region, access rights
- Several of these are created for each mapped image
 - Executable code, initialized data, uninitialized data

Demand Paging: Page Fault Handling

- Soon the process will access an address without a valid page
 - OS gets a page fault from the MMU
- What happens?
 - Kernel searches a tree structure of memory allocations for the process to see if the faulting address is valid
 - If not valid, send a SEGV signal to the process
 - Is the type of access valid for the page?
 - Send a signal if not
 - We have a valid page but it's not in memory

Keeping track of a processes' memory region

Linux stores information about a process' address space in the **memory descriptor** of the PCB.



Demand Paging: Getting a Page

- The page we need is either in the a mapped file (executable or library) or in a swap file
 - If PTE is not valid but page # is present
 - The page we want has been saved to a swap file
 - Page # in the PTE tells us the location in the file
 - If the PTE is not valid and no page #
 - Load the page from the program file from the disk
- Read page into physical memory
 - Find a free page frame (evict one if necessary)
 - Read the page: This takes time: context switch & block
 - Update page table for the process
 - Restart the process at the instruction that faulted

Page Replacement

- A process can run without having all of its memory allocated
 - It's allocated on demand
- If the
 $\{\text{address space used by all processes} + \text{OS}\} \leq \text{physical memory}$
then we're ok
- Otherwise:
 - Make room: discard or store a page onto the disk
 - If the page came from a file & was not modified
 - Discard ... we can always get it
 - If the page is dirty, it must be saved in a **swap file**
 - Swap file: a file (or disk partition) that holds excess pages

Cost

- Handle page fault exception: ~ 400 usec
- Disk seek & read: ~ 10 msec
- Memory access: ~ 100 ns
- Page fault degrades performance by around 100,000!!
- Avoid page faults!
 - If we want < 10% degradation of performance, we can have just one page fault per 1,000,000 memory accesses

Page replacement

- We need a good replacement policy for good performance

FIFO Replacement

- First In, First Out
- Good
 - May get rid of initialization code or other code that's no longer used
- Bad
 - May get rid of a page holding frequently used global variables

Least Recently Used (LRU)

- Timestamp a page when it is accessed
- When we need to remove a page, search for the one with the oldest timestamp
- Nice algorithm but...
 - Timestamping is a pain – we can't do it with the MMU!

Not Frequently Used Replacement

- Approximate LRU
- Each PTE has a reference bit
- Keep a counter for each page frame
- At each clock interrupt:
 - Add the reference bit of each frame to its counter
 - Clear reference bit
- To evict a page, choose the frame with the lowest counter
- Problem
 - No sense of time: a page that was used a lot a long time ago may still have a high count
 - Updating counters is expensive

Clock (Second Chance)

- Arrange physical pages in a logical circle (circular queue)
 - Clock hand points to first frame
- Paging hardware keeps 1 *reference* bit per frame
 - Set *reference* bit on memory reference
 - If it's not set then the frame hasn't been used for a while
- On page fault:
 - Advance clock hand
 - Check *reference* bit
 - If 1, it's been used recently – clear & advance
 - If 0, evict this page

Enhanced Clock (Second Chance)

- Use the *reference* and *modify* bits of the page
- Choices for replacement – (reference, modify):
 - (0, 0): not referenced recently or modified
 - Good candidate for replacement
 - (0, 1): not referenced recently but modified.
 - The page will have to be saved before replacement
 - (1, 0): recently used.
 - Less ideal – will probably be used again
 - (1, 1): recently used and modified
 - Least ideal – will probably be used again AND we'll have to save it to a swap file if we replace it.
- Algorithm: like clock but replace the first page in the lowest non-empty class

Nth Chance Replacement

- Similar to Second Chance
- Maintain a counter along with a reference bit
- On page fault:
 - Advance clock hand
 - Check reference bit
 - If 1, clear and set counter to 0
 - If 0, increment counter. If counter $< N$, go on. Else evict
- Better approximation of LRU

Kernel Swap Daemon

- *kswapd* on Linux
- Anticipate problems
- Decides whether to shrink caches if page count is low
 - Page cache, buffer cache
 - Evict pages from page frames

Demand paging summary

- Allocate page table
 - Map kernel memory
 - Initialize stack
 - Memory-map text & data from executable program (& libraries)
 - But don't load!
- Load pages on demand (first access)
 - When we get a page fault

Summary: If we run out of free page frames

- Free some page frames
 - Discard pages that are mapped to a file
or
 - Move some pages to a **swap** file
- Clock algorithm
- Anticipate need for free page frames
 - *kswapd* – kernel swap daemon

Multitasking Considerations

Supporting multitasking


- Multiple address spaces can be loaded in memory
- A CPU register points to the current page table
- OS changes the register set when context switching
- Performance increased with Address Space ID in TLB

Working Set

- Keep active pages in memory
- A process needs its working set in memory to perform well
 - **Working set** = set of pages that have been referenced in the last window of time
 - **Spatial locality**
 - Size of working set varies during execution
- More processes in a system:
 - Good: increase throughput; chance that some process is available to run
 - Bad: **thrashing**: processes do not have enough page frames available to run without paging

Thrashing

- Locality
 - Process migrates from one locality (working set) to another
- Thrashing
 - Occurs when sum of all working sets $>$ total memory



Resident Set Management

- Resident set = set of a process' pages in memory
- How many pages of a process do we bring in?
- Resident set can be **fixed** or **variable**
- Replacement scope: **global** or **local**
 - Global: process can pick a replacement from *all* frames
- Variable allocation with global scope:
 - Simple
 - Replacement policy may not take working sets into consideration
- Variable allocation with local scope
 - More complex
 - Modify resident size to approximate working set size

Working Set Model

- Approximates locality of a program
- Δ : *working set window*:
 - Amount of elapsed time while the process was actually executing (e.g., count of memory references)
- WSS_i : working set size of process P_i
 - WSS_i = set of pages in most recent Δ page references
- System-wide demand for frames
$$D = \sum WSS_i$$
- If $D > \text{total memory size}$, then we get thrashing

Page fault frequency

- Too small a working set causes a process to thrash
- Monitor page fault frequency per process
 - If too high, the process needs more frames
 - If too low, the process may have too many frames

Dealing with thrashing

- If all else fails ...
- Suspend a process(es)
 - Lowest priority, Last activated, smallest resident set, ...?
- Swapping:
 - Move an entire process onto the disk: no pages in memory
 - Process must be re-loaded to run

The End