

Operating Systems Design

13. File System Implementation

Paul Krzyzanowski
pxk@cs.rutgers.edu

Terms

- **Disk**
 - Non-volatile block-addressable storage.
- **Block = sector**
 - Smallest chunk of I/O on a disk
 - Most disks have 512-byte blocks
- **Partition**
 - Subset of all blocks on a disk. A disk has ≥ 1 partitions
- **Volume**
 - Disk, disks, or partition that contains a file system
 - A volume may span disks

More terms

- **Superblock**
 - Area on the volume that contains key file system information
- **Metadata**
 - Attributes of a file, not the file contents (data)
 - E.g., modification time, length, permissions, owner
- **inode**
 - A structure that stores a file's metadata and location of its data

Files

- **Contents (Data)**
 - Unstructured (byte stream) or structured (records)
 - Stored in data blocks
 - Find a way of allocating and tracking the blocks that a file uses
- **Metadata**
 - Usually stored in an inode ... sometimes in a directory entry
 - Except for the name, which is stored in a directory

Directories

- A directory is just a file containing names & references
 - Name \rightarrow (metadata, data) *Unix (UFS) approach*
 - (Name, metadata) \rightarrow data *MS-DOS (FAT) approach*
- **Linear list**
 - Search can be slow for large directories.
 - Cache frequently-used entries
- **Hash table**
 - Linear list but with hash structure
 - Hash(name)
- More exotic structures: **B-Tree, HTree**

Block allocation: **Contiguous**


- Each file occupies a set of adjacent blocks
- You just need to know the starting block & file length
- We'd love to have contiguous storage for files!
 - Minimize disk seeks when accessing a file

Problems with contiguous allocation

- Storage allocation is a pain (remember main memory?)
 - **External fragmentation**: free blocks of space scattered throughout
 - vs. **Internal fragmentation**: unused space within a block (allocation unit)
 - Periodic defragmentation: move files (yuck!)
- Concurrent file creation: how much space do you need?
- Compromise solution: **extents**
 - Allocate a contiguous chunk of space
 - If the file needs more space, allocate another chunk (extent)
 - Need to keep track of all extents
 - Not all extents will be the same size: it depends how much contiguous space you can allocate

Block allocation: **Linked Allocation**

- A file's data is a linked list of disk blocks
 - Directory contains a pointer to the first block of the file
 - Each block contains a pointer to the next block
- Problems
 - Only good for sequential access
 - Each block uses space for the pointer to the next block
- **Clusters**
 - Multiples of blocks: reduce overhead for block pointer & improve throughput
 - *A cluster is the smallest amount of disk space that can be allocated to a file*
 - Penalty: increased **internal fragmentation**



File Allocation Table (DOS/Windows FAT)

- Variation of Linked Allocation
- Section of disk at beginning of the volume contains a file allocation table
- The table has one entry per block. Contents contain the next logical block (cluster) in the file.


Directory entry:

myfile.txt	metadata	06
------------	----------	----

FAT table: one per file system

0	0	0	12	0	0	03	0	0	0	0	0	-1	0
---	---	---	----	---	---	----	---	---	---	---	---	----	---

- FAT-16: 16-bit block pointers
 - 16-bit cluster numbers; up to 64 sectors/cluster
 - Max file system size = 2 GB (with 512 byte sectors)
- FAT-32: 32-bit block pointers
 - 32-bit cluster numbers; up to 64 sectors/cluster
 - Max file system size = 8 TB (with 512 byte sectors)
 - Max file size = 4 GB

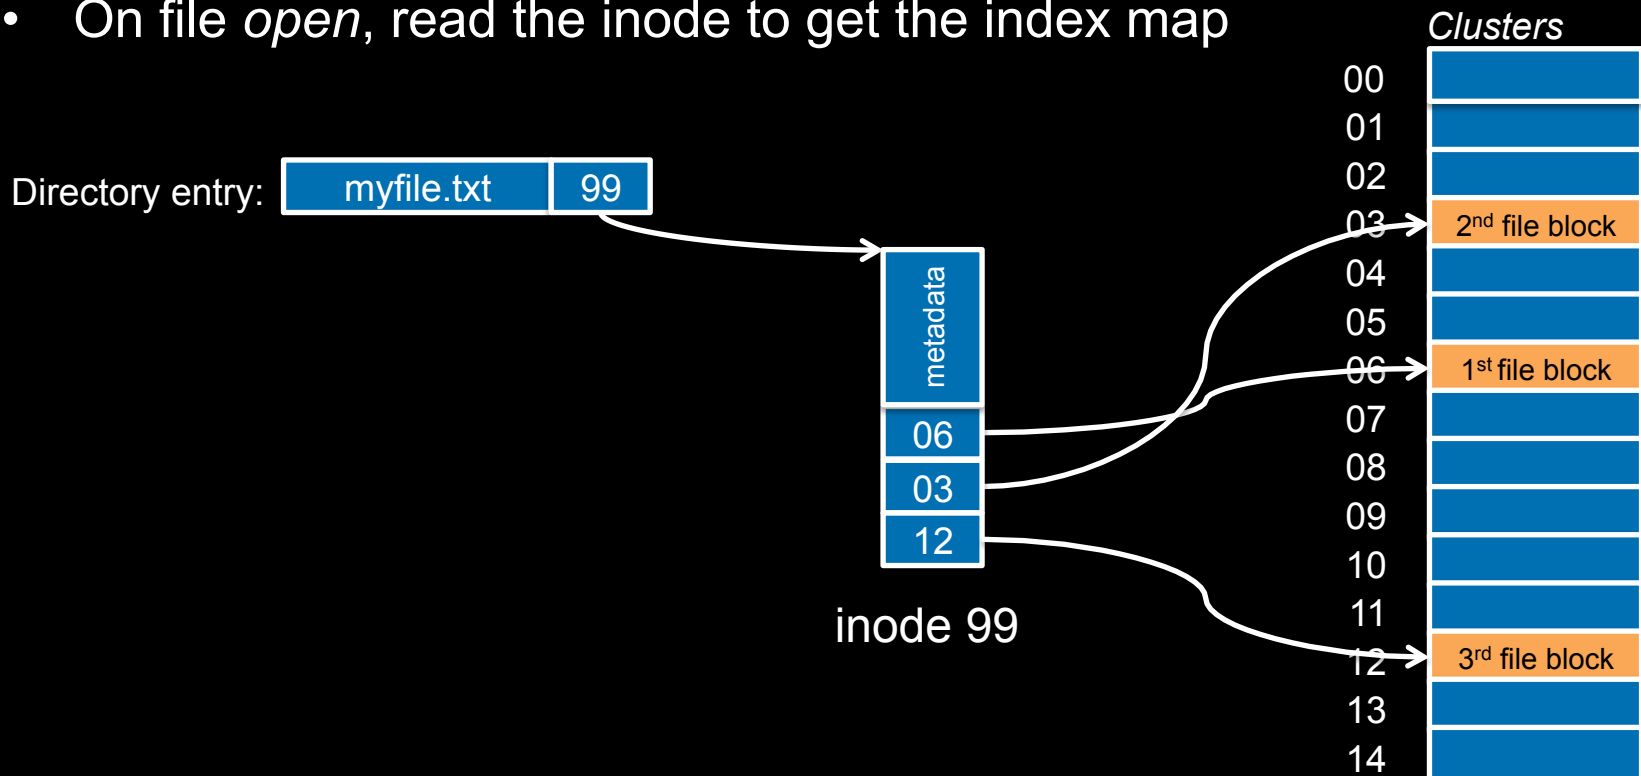


Indexed Allocation (Block map)

- Linked allocation is not efficient for random access
- FAT requires storing the *entire* table in memory for efficient access
- **Indexed allocation:**
 - Store the entire list of block pointers for a file in one place: the index block (**inode**)
 - One inode per file
 - We can read this into memory when we open the file

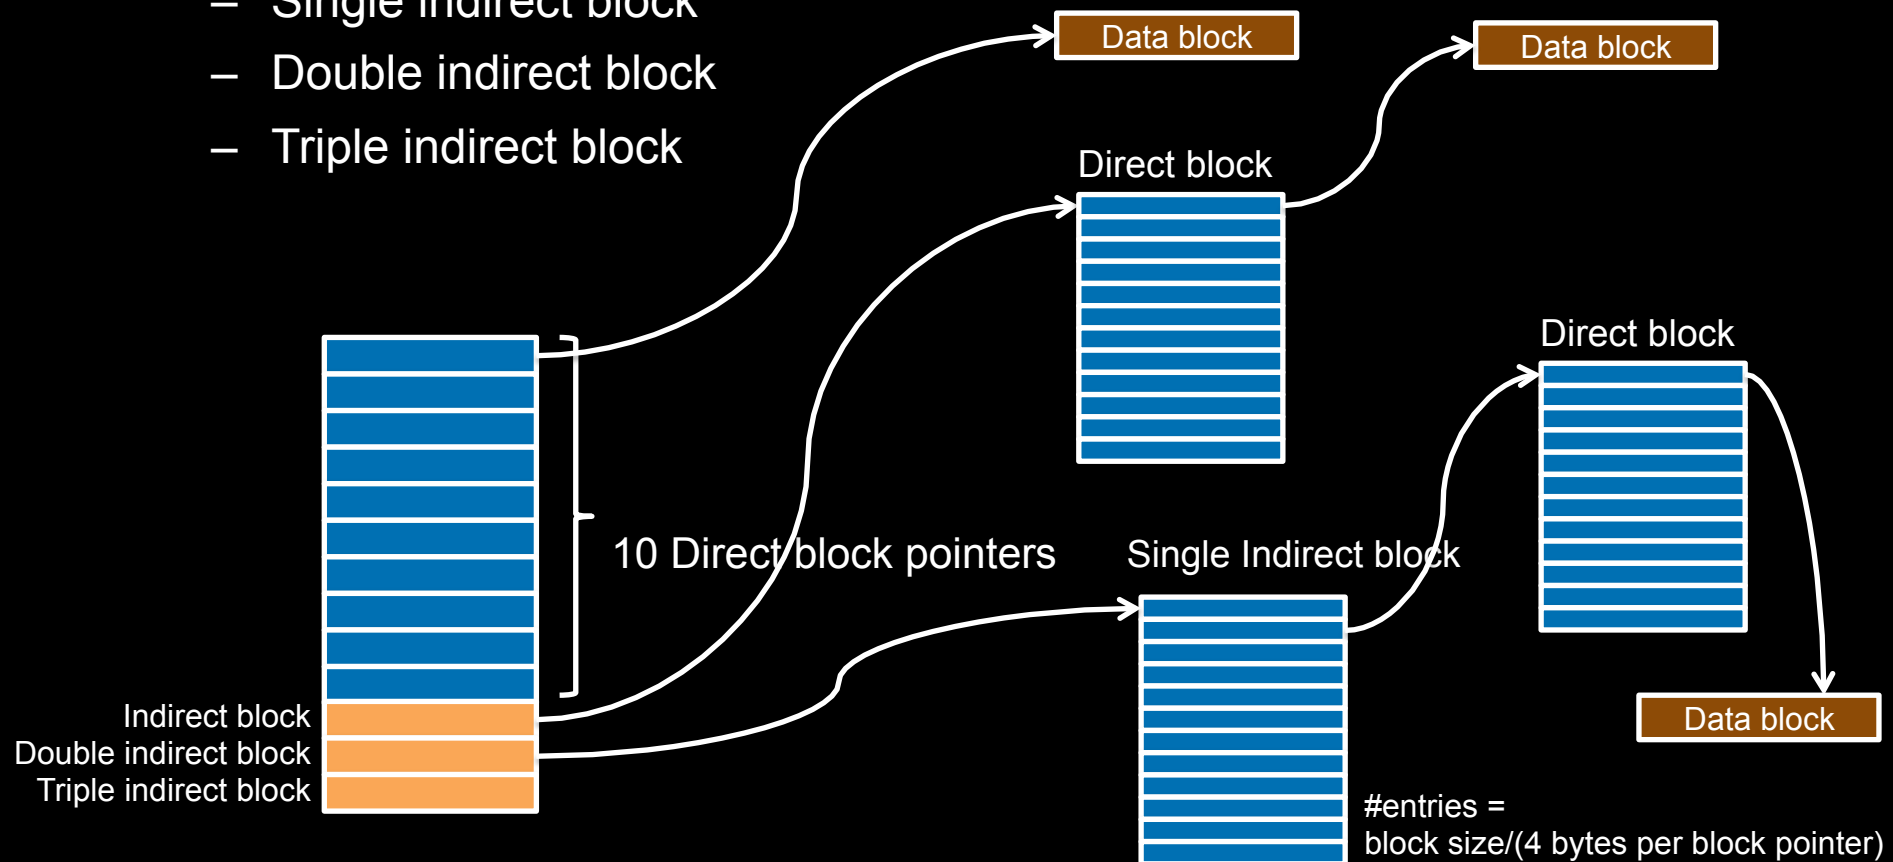
Indexed Allocation (block/cluster map)

- Directory entry contains name and inode number
- inode contains file metadata (length, timestamps, owner, etc.) *and* a block map
- On file *open*, read the inode to get the index map



Combined indexing (Unix File System)

- We want inodes to be a fixed size
- Large files get
 - Single indirect block
 - Double indirect block
 - Triple indirect block



Combined Indexing: inside the inode

- **Direct block numbers**
 - These contain block numbers that contain the file's data
- **Indirect block number**
 - This is a block number of a block that contains a list of direct block numbers. Each block number is the number of a block that contains the file's data
- **Double indirect block number**
 - This refers to a block that contains a list of indirect block numbers. Each indirect block number is the number of a block that contains a list of direct block numbers
- **Triple indirect block number**
 - This refers to a block that contains a list of double indirect block numbers. Each double indirect block number is the number of a block that contains a list of indirect direct block numbers. Each of these contains a list of direct block numbers

Example

- Unix File System
 - 1024-byte blocks, 32-bit block pointers
 - inode contains
 - 10 direct blocks, 1 indirect, 1 double-indirect, 1 triple indirect
- Capacity
 - Direct blocks will address: $1K \times 10 \text{ blocks} = 10,240 \text{ bytes}$
 - 1 level of indirect blocks: additional $(1K/4) \times 1K = 256K \text{ bytes}$
 - 1 Double indirect blocks: additional $(1K/4) \times (1K/4) \times 1K = 64M \text{ bytes}$
 - Maximum file size = $10,240 + 256K + 64M = 65,792K \text{ bytes}$

Extent lists

- **Extents:** Instead of listing block addresses
 - Each address represents a range of blocks
 - Contiguous set of blocks
 - E.g., 48-bit block # + 2-byte length (total = 64 bits)
- Why are they attractive?
 - Less block numbers to store if we have lots of contiguous allocation
- Problem: file seek operations
 - Locating a specific location requires traversing a list
 - Extra painful with indirect blocks

Directories

- A directory is just a file
 - Name \rightarrow (metadata, data)
 - or (Name, metadata) \rightarrow data
- **Linear list**
 - Search can be slow for large directories
 - Cache frequently-used entries
- **Hash table**
 - Linear list but with hash structure
 - Hash(name)
- **B-Tree**
 - Balanced tree (constant depth) with high fan-out
 - Variations include B+ Tree and HTree

Implementing File Operations

Initialization

- **Low-level formatting** (file system independent)
 - Define blocks (sectors) on a track
 - Create spare sectors
 - Identify and remap bad blocks
- **High-level formatting** (file system specific)
 - Define the file system structure
 - Initialize the free block map
 - Initialize sizes of inode and journal areas
 - Create a top-level (root) directory

Opening files

Two-step process

1. Pathname **Lookup** (*namei* function in kernel)
 - Traverse directory structure based on the pathname to find file
 - Return the associated inode
 - (cache frequently-used directory entries)

2. **Verify** access permissions
 - If OK, allocate in-memory structure to maintain state about access (e.g., that file is open read-only)

Writing files

A *write* either overwrites data in a file or adds data to the file, causing it to grow

- Allocate disk blocks to hold data
- Add the blocks to the list of blocks owned by the file
 - Update free block bitmap, the inode, and possibly indirect blocks
 - Write the file data
 - Modify inode metadata (file length)
 - Change current file offset in kernel

Deleting files

- Remove name from the directory
 - Prevent future access
- If there are no more links to the inode (disk references)
 - mark the file for deletion
- ... and if there are no more programs with open handles to the file (in-memory references)
 - Release the resources used by the file
 - Return data blocks to the free block map
 - Return inode to the free inode list
- Example:
 - Open temp file, delete it, continue to access it
 - OS cleans up the data when the process exits

Additional file system operations

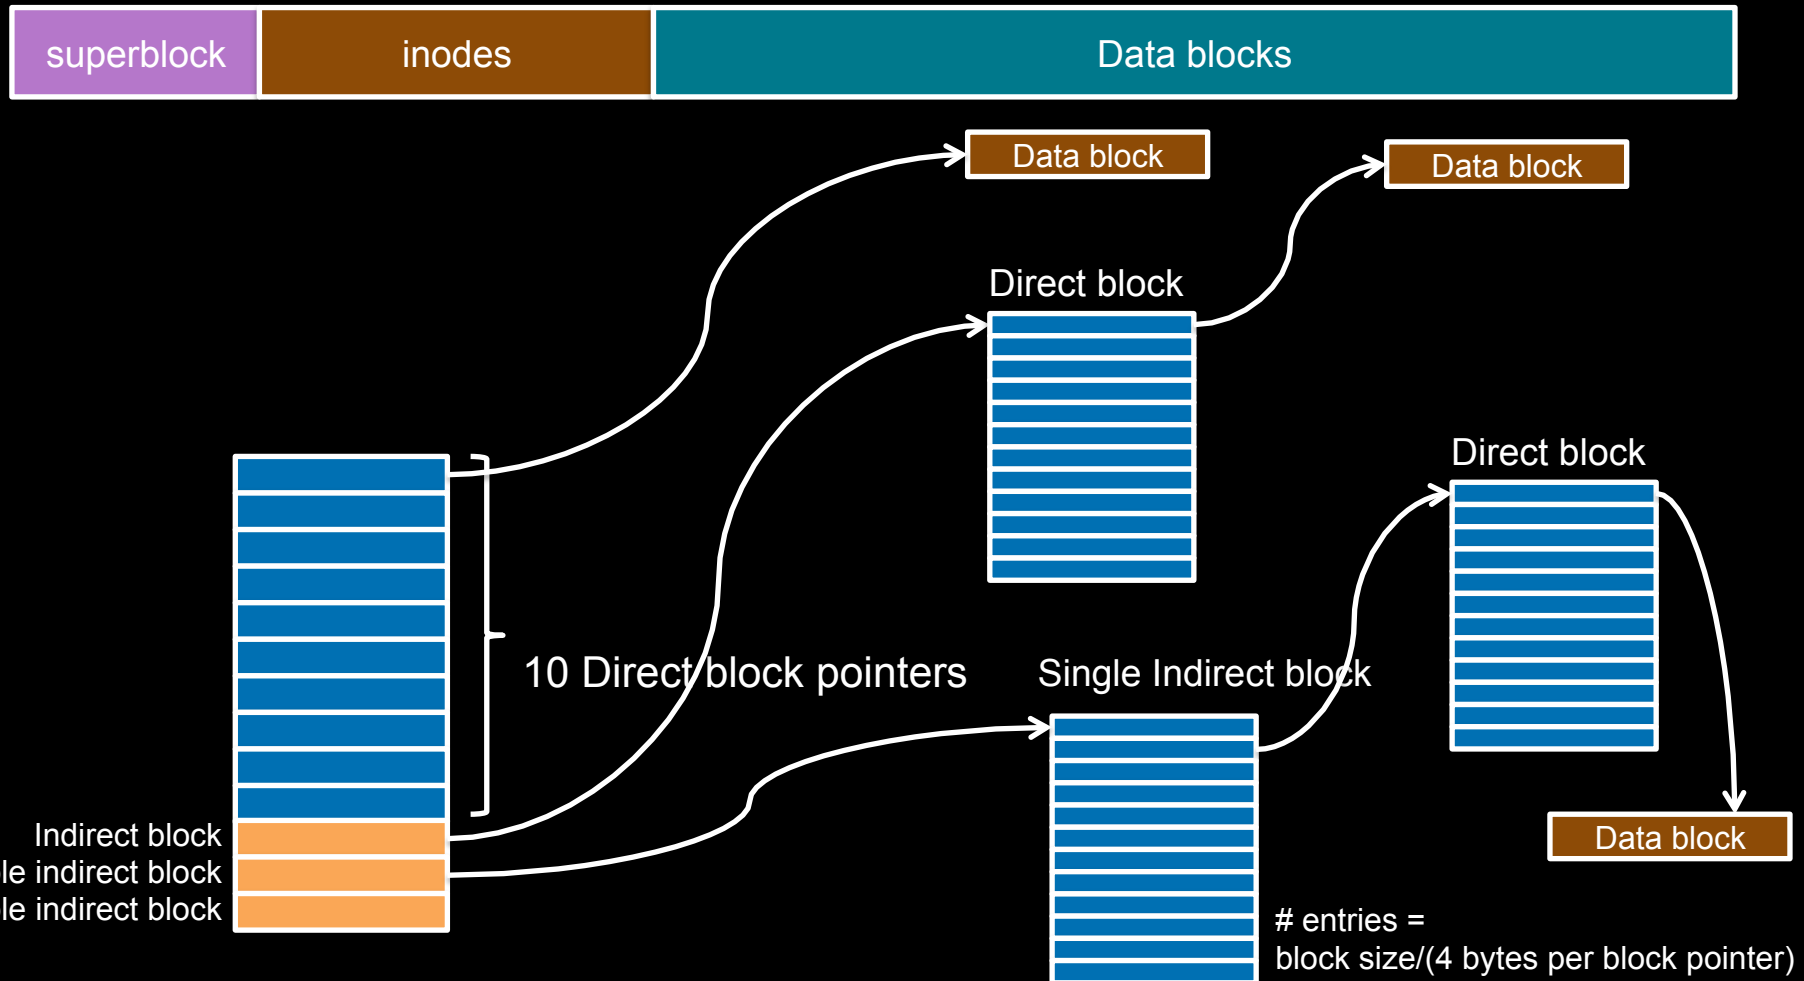
- **Hard links** (aliases)
 - Multiple directory entries (file names) that refer to the same inode
 - inode contains reference count to handle deletion
- **Symbolic links**
 - File data contains a path name
 - Underlying file can disappear or change
- **Access control lists** (ACLs)
 - Classic UNIX approach: user, group, world permissions
 - ACL: enumerated list of users and permissions
 - Variable size

Additional file system operations

- **Extended attributes** (NTFS, HFS+, XFS, etc.)
 - E.g., store URL from downloaded web/ftp content, app creator, icons
- **Indexing**
 - Create a database for fast file searches
- **Journaling**
 - Batch groups of changes. Commit them at once to a transaction log

Unix File System (UFS)

inodes with direct, indirect, double-indirect, and triple-indirect blocks



Unix File System (UFS)

Superblock contains:

- Size of file system
- # of free blocks
- list of free blocks (+ pointer to free block lists)
- index of the next free block in the free block list
- Size of the inode list
- Number of free inodes in the file system
- Index of the next free inode in the free inode list
- Modified flag (clean/dirty)

Unix File System (UFS)

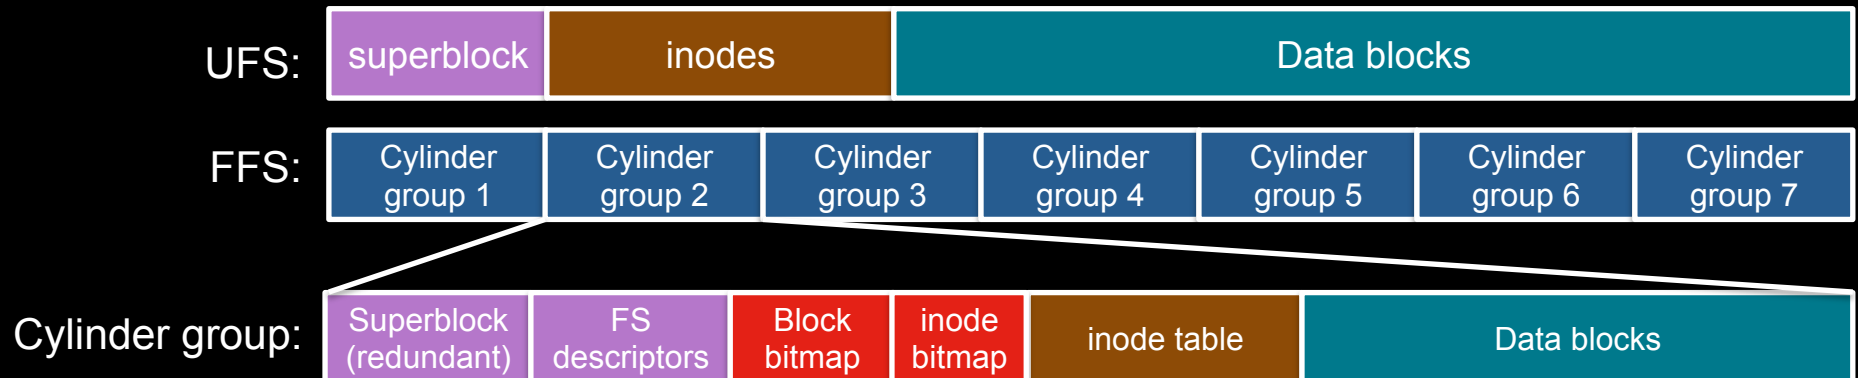
- Free space managed as a linked list of blocks
 - Eventually this list becomes random
 - Every disk block access will require a seek!
- Fragmentation is a big problem
- Typical performance was often:
2–4% of raw disk bandwidth!

BSD Fast File System (FFS)

- Try to improve UFS
- **Improvement #1: Use larger blocks**
 - ≥ 4096 bytes instead of UFS's 512-byte or 1024-byte blocks
 - Block size is recorded in the superblock
 - Just doubling the block size resulted in $> 2x$ performance!
 - 4 KB blocks let you have 4 GB files with only two levels of indirection
 - Problem: increased internal fragmentation
 - Lots of files were small
 - Solution: Manage fragments within a block (down to 512 bytes)
 - A file is 0 or more full blocks and possibly one fragmented block
 - Free space bitmap stores fragment data
 - As a file grows, fragments are copied to larger fragments and then to a full block
 - Allow user programs to find the optimal block size
 - Standard I/O library and others use this
 - Also, avoid extra writes by caching in the system buffer cache

BSD Fast File System (FFS)

- **Improvement #2: Minimize head movement (reduce seek time)**
 - Seek latency is usually much higher than rotational latency
 - Keep file data close to its inode to minimize seek time to fetch data
 - Keep related files & directories together
 - Cylinder: collection of all blocks on the same track on all heads of a disk
 - **Cylinder group**: Collection of blocks on one or more consecutive cylinders



How do you find inodes?

- UFS was easy:
 - $\text{inodes_per_block} = \text{sizeof}(\text{block}) / \text{sizeof}(\text{inode})$
 - $\text{inode_block} = \text{inode} / \text{inodes_per_block}$
 - $\text{block_offset} = (\text{inode} \% \text{inodes_per_block}) * \text{sizeof}(\text{inode})$
- FFS
 - We need to know how big each chunk of inodes in a cylinder group is: keep a table

BSD Fast File System (FFS)

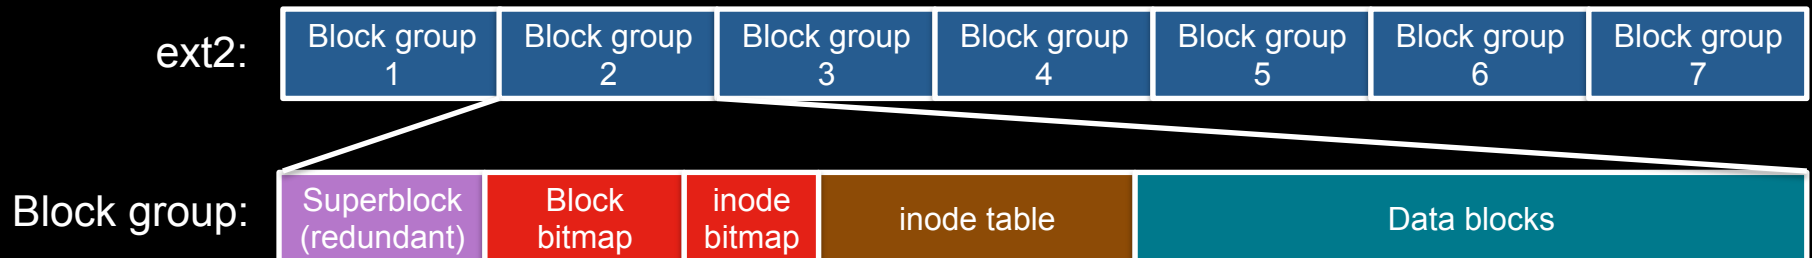
- Optimize for sequential access
- Allocate data close together
 - Pre-allocate up to 8 adjacent blocks when allocating a block
 - Achieves good performance under heavy loads
 - Speeds sequential reads
- Prefetch
 - If 2 or more logically sequential blocks are read
 - Assume sequential read and request one large I/O on the entire range of sequential blocks
 - Otherwise, schedule a read-ahead

BSD Fast File System (FFS)

- **Improve fault tolerance**
 - Strict ordering of writes of file system metadata
 - *fsck* still requires up to five passes to repair
 - All metadata writes are synchronous (not buffered)
 - This limits the max # of I/O operations
- Directories
 - Max filename length = 256 bytes (vs. 12 bytes of UFS)
- **Symbolic links** introduced
 - Hard links could not point to directories and worked only within the FS
- Performance:
 - 14-47% of raw disk bandwidth
 - Better than the 2-5% of UFS

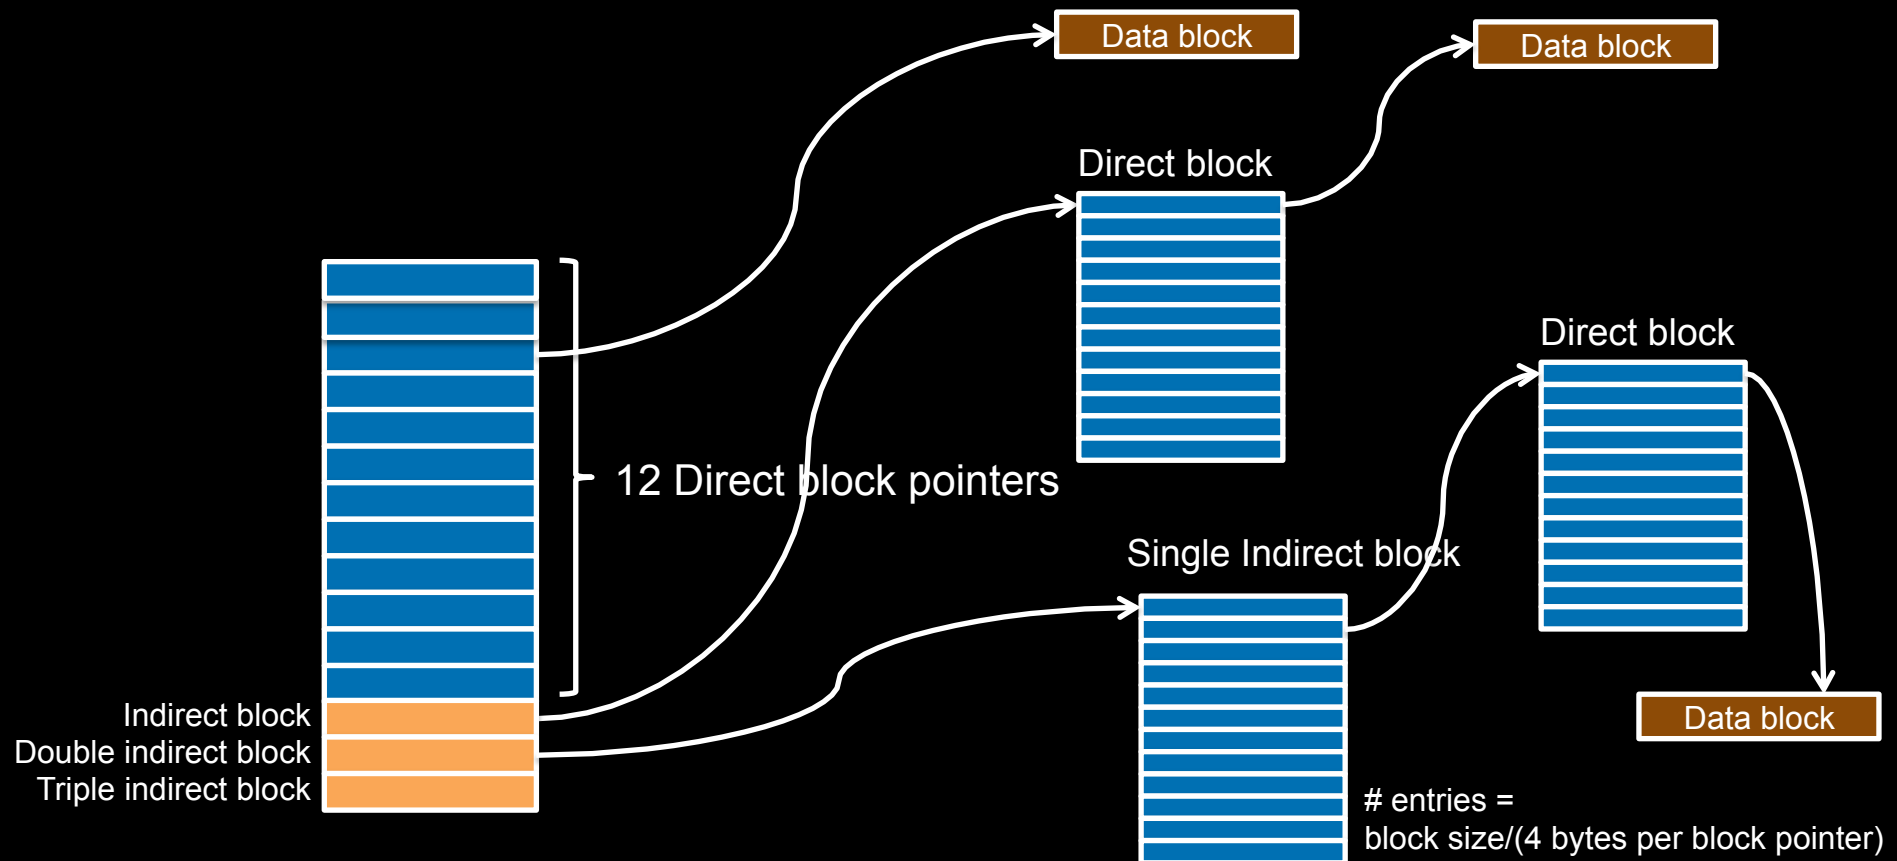
Linux ext2

- Similar to BSD FFS
- No fragments
- No cylinder groups (not useful in modern disks) – block groups
- Divides disk into fixed-size block groups
 - Like FFS, somewhat fault tolerant: recover chunks of disk even if some parts are not accessible



Linux ext2

inodes with direct, indirect, double-indirect, and triple-indirect blocks



Linux ext2


- Improve performance via aggressive caching
 - Reduce fault tolerance because of no synchronous writes
 - Almost all operations are done in memory until the buffer cache gets flushed
- Unlike FFS:
 - No guarantees about the consistency of the file system
 - Don't know the order of operations to the disk: risky if they don't all complete
 - No guarantee on whether a write was written to the disk when a system call completes
- In most cases, ext2 is *much* faster than FFS

Journaling

File system inconsistencies

Example:

- Writing a block to a file may require:
 - inode is
 - updated with a new block pointer
 - Updated with a new file size
 - Data free block bitmap is updated
 - Data block contents written to disk
- If all of these are not written, we have a file system inconsistency
- **Consistent update problem**



Journaling

- Journaling = write-ahead logging
- Keep a transaction-oriented journal of changes
 - Record what you are about to do (*along with the data*)

```
Transaction-begin  
  New inode 779  
  New block bitmap, group 4  
  New data block 24120  
Transaction-end
```

- Once this has committed to the disk then overwrite the real data
 - If all goes well, we don't need this transaction entry
 - If a crash happens any time after the log was committed
 Replay the log on reboot (redo logging)
- This is called *full data journaling*

Writing the journal

- Writing the journal all at once would be great but is risky
 - We don't know what order the disk will schedule the block writes
 - Don't want to risk having a "transaction-end" written while the contents of the transaction have not been written yet
 - Write all blocks *except* transaction-end
 - Then write transaction-end
- If the log is replayed and a transaction-end is missing, ignore the log entry

Cost of journaling

- We're writing everything twice
 - ...and constantly seeking to the journal area of the disk
- Optimization
 - Do not write user data to the journal
 - **Metadata journaling** (also called **ordered journaling**)

```
Transaction-begin
  New inode 779
  New block bitmap, group 4
Transaction-end
```

- What about the data?
 - Write it to the disk first (not in the journal)
 - Then mark the end of the transaction
 - This prevents pointing to garbage after a crash and journal replay

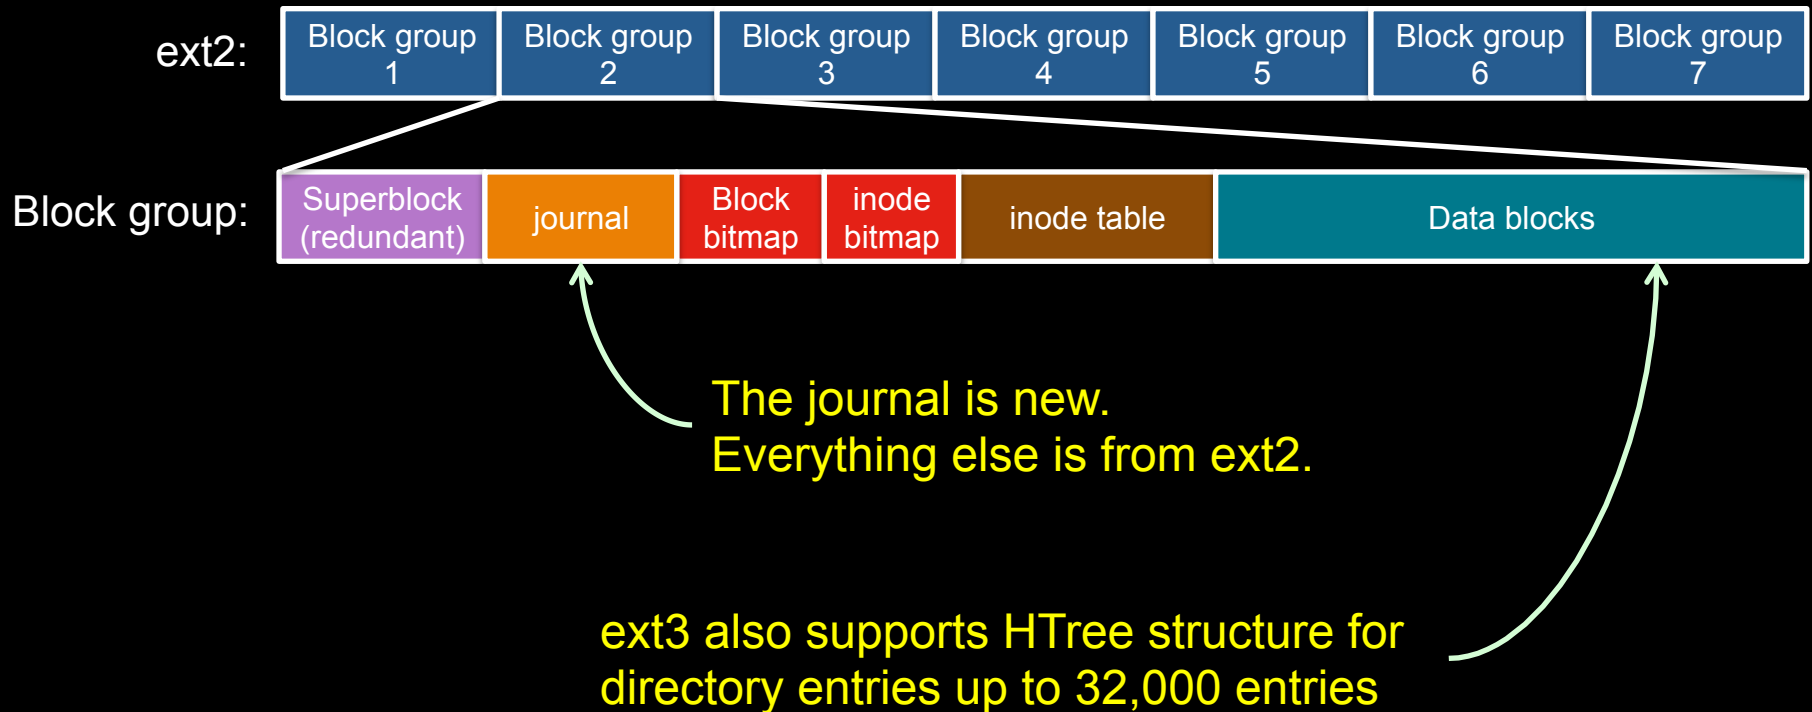
Linux ext3

- ext3 = ext2 + **journaling** (mostly)
- Goal: improved fault recovery
 - Reduce the time spent in checking file system consistency & repairing the file system

ext3 journaling options

- **journal**
 - full data + metadata journaling
 - [slowest]
- **ordered**
 - Data blocks written first, then metadata journaling
 - Write a transaction-end only when the other writes have completed
- **writeback**
 - Metadata journaling with no ordering of data blocks
 - Recent files can get corrupted after a crash
 - [fastest]

ext3 layout



Linux ext4: extensions to ext3


- Large file system support
 - 1 exabyte (10^{18} bytes); file sizes to 16 TB
- **Extents** used instead of block maps
 - Range of contiguous blocks
 - 1 extent can map up to 12 MB of space (4 KB block size)
 - 4 extents per inode. Additional ones are stored in an HTree (constant-depth tree similar to a B-tree)
- Ability to pre-allocate space for files
 - Increase chance that it will be contiguous
- Delayed allocation
 - Allocate on flush – only when data is written to disk
 - Improve block allocation decisions because we know the size

Linux ext4: extensions to ext3

- Over 64,000 directory entries (vs. 32,000 in ext3)
 - HTree structure
- Journal checksums
 - Monitor journal corruption
- Faster file system checking
 - Ignore unallocated block groups
- Interface for multiple-block allocations
 - Increase contiguous storage
- Timestamps in nanoseconds

Microsoft NTFS

- Standard file system for Windows; successor to FAT-32
- 64-bit volume sizes, journaling, and data compression
- Cluster-based (file compression not supported on clusters > 4 KB)



Boot Sector: info about layout of the volume & FS structures; Windows bootloader

MFT: contains information about all files in the file system

File system data: all the data that is not in the MFT

MFT Copy: copy of critical part of MFT for recovery (first 4 records)

NTFS Master File Table

- The MFT is itself a file (starting at a well-known place)
- It contains file records (inode) for all files, including itself
 - B-Tree structure
- MFT Special files:

MFT record 0	\$Mft	Master file table
MFT record 1	\$MftMirr	Duplicate of 1 st 4 records of MFT
MFT record 2	\$LogFile	Metadata journal for recovery
MFT record 3	\$Volume	Info about the file system volume
MFT record 4	\$AttrDef	Attribute definitions
MFT record 5		Root folder
MFT record 6	\$Bitmap	Cluster bitmap (free/used clusters)

And a few more less interesting ones...

- Because the Bitmap is just a file, the volume bitmap is a file, the size of a volume can be easily expanded

NTFS MFT & Attributes

- MFT can grow just like any other file
 - To minimize fragmentation, 12.5% of the volume is reserved for use by the MFT (“MFT Zone”)
- Each file record is 1, 2, or 4 KB (determined at FS initialization)
- File record info: set of typed attributes
 - Some attributes may have multiple instances (e.g., name & MS-DOS name)
 - **Resident attributes**: attributes that fit in the MFT record
 - If the attributes take up too much space, additional clusters are allocated
 - an “**Attribute List**” attribute is added
 - Describes location of all other file records
 - Attributes stored outside of the MFT record are **Nonresident attributes**

NTFS File Data

- File data is an attribute
 - NTFS supports multiple data attributes per file
 - One main, unnamed stream associated with a data file; other named streams are possible
 - Manage related data as a single unit
- Small folders and small data files can fit entirely within the MFT.
 - Large folders are B-tree structures and point to external clusters
- Block allocation: via extents

Microsoft NTFS

- Directories
 - Stored as B+ trees in alphabetic order
 - Name, MFT location, size of file, last access & modification times
 - Size & times are duplicated in the file record & directory entry
 - Designed to optimize some directory listings
- Write-ahead logging
 - Writes planned changes to the log, then writes the blocks
- Transparent data compression of files
 - Method 1: compress long ranges of zero-filled data by not allocating them to blocks (sparse files)
 - Method 2: break file into 16-block chunks
 - Compress each chunk
 - If at least one block is not saved then do not compress the chunk

The next generation: ReFS for Windows 8

- ReFS = Resilient File System
- Goals
 - Verify & auto-correct data; checksums for metadata
 - Optimize for extreme scale
 - Never take the file system offline – even in case of corruption
 - Allocate-on-write transactional model
 - Shared storage pools for fault tolerance & load balancing
 - Data striping for performance; redundancy for fault tolerance
- General approach
 - Use B+ trees to represent all information on the disk
 - “Table” interface for enumerable sets of key-value pairs
 - Provide a generic key-value interface to implement files, directories, and all other structures

Log Structured File Systems

NAND flash memory

- Memory arranged in “pages” – similar to disk sectors
 - Unit of allocation and programming
 - Individual bytes cannot be written
- Conventional file systems
 - Modify the same blocks over and over
 - At odds with NAND flash performance
 - Also, optimizations for seek time are useless
- Limited erase-write cycles
 - 100,000 to 1,000,000 cycles
 - Employ wear leveling to distribute writes among all (most) blocks
 - Bad block “retirement”
- Options:
 - Managed NAND = NAND flash + integrated controller
 - Handles block mapping
 - Can use conventional file systems
 - OS file system software optimized for Flash

Dynamic wear leveling

- **Dynamic wear leveling**

- Monitors high- and low-use areas of the memory
 - At a certain point it will swap high-use blocks with low-use ones
- Map logical block addresses to flash memory block addresses
- When a block of data is written to flash memory, it is written to a new location and the map is updated
- Blocks that are never modified will not get reused

- **Static wear leveling**

- Static data is moved periodically to give all blocks a chance to wear evenly

File systems designed for wear leveling

- **YAFFS2** and **JFFS2**
 - YAFFS is favored for disks > 64 MB
 - File system of choice for Android
 - JFFS is favored for smaller disks
 - Typically used in embedded systems
 - Only address dynamic wear leveling

YAFFS

- Unit of allocation = “chunk”
- Several (32 – 128+) chunks = 1 block
 - Unit of erasure for YAFFS
- Log structure: all updates written sequentially
 - Each log entry is 1 chunk in size:
 - *Data chunk*
 - or *Object header* (describes directory, file, link, etc.)
 - Sequence numbers are used to organize a log chronologically
 - Each chunk contains:
 - Object ID: object the chunk belongs to
 - Chunk ID: where the chunk belongs in the file
 - Byte count: # bytes of valid data in the chunk

YAFFS

Create a file

Chunk	ObjectId	ChunkID		
0	500	0	Live	Object header for file (length=0)
1				
2				
3				

Block 1

Adapted from <http://www.yaffs.net/files/yaffs.net/HowYaffsWorks.pdf>

YAFFS

Write some data to the file

Chunk	ObjectId	ChunkID		
0	500	0	Live	Object header for file (length=0)
1	500	1	Live	First chunk of data
2	500	2	Live	Second chunk of data
3	500	3	Live	Third chunk of data

Block 1

Adapted from <http://www.yaffs.net/files/yaffs.net/HowYaffsWorks.pdf>

YAFFS

Close the file: write new header

Chunk	ObjectId	ChunkID			
0	500	0	Deleted	Object header for file (length=0)	Block 1
1	500	1	Live	First chunk of data	
2	500	2	Live	Second chunk of data	
3	500	3	Live	Third chunk of data	
0	500	0	Live	Object header for file (length= n)	Block 2

Adapted from <http://www.yaffs.net/files/yaffs.net/HowYaffsWorks.pdf>

YAFFS

Open file; modify first chunk; close file

Chunk	ObjectId	ChunkID			
0	500	0	Deleted	Object header for file (length=0)	Block 1
1	500	1	Deleted	First chunk of data	
2	500	2	Live	Second chunk of data	
3	500	3	Live	Third chunk of data	
0	500	0	Deleted	Object header for file (length= n)	Block 2
1	500	1	Live	New first chunk of data	
2	500	0	Live	New object header for file (length= n)	

Adapted from <http://www.yaffs.net/files/yaffs.net/HowYaffsWorks.pdf>

YAFFS Garbage Collection

- If all chunks in a block are deleted, the block can be erased & reused
- If blocks have some free chunks
 - We need to do **garbage collection**
 - Copy active chunks onto other blocks so we can free a block
 - **Passive collection**: pick blocks with few used chunks
 - **Aggressive collection**: try harder to consolidate chunks

YAFFS in-memory structures

- Construct file system state in memory
 - Map of in-use chunks
 - In-memory object state for each object
 - File tree/directory structure to locate objects
 - Scan the log backwards chronologically
highest→lowest sequence number
 - Checkpointing: save the state of these structures at unmount time to speed up the next mount

YAFFS error detection/correction

- ECC used for error recovery
 - Correct 1 bad bit per 256 bytes
 - Detect 2 bad bits per 256 bytes
 - Bad blocks:
if read or write fails, ask driver to mark the block as bad

JFFS1

- **Log-structured file system**
 - Nodes containing data & metadata are stored sequentially, progressing through the storage space: **circular log**
 - Node contains inode number, all metadata, and variable data
 - Each new node contains a version higher than a previous one
 - inode numbers are 32-bits; never reused
 - Old nodes (for data that is covered in later nodes) are “dirty space”
- On mount, the entire file system is scanned to rebuild the directory hierarchy
- Garbage collection: go through dirty space – find unneeded nodes and data blocks that are overwritten

JFFS2

- Keep linked lists representing individual erase blocks
 - Most of erase blocks will be on the *clean list* or *dirty list*
- Roughly 99/100 of the time, pick a block from the dirty list. Otherwise, pick a block from the clean list.
 - Optimize garbage collection to reuse blocks which are partially obsolete
 - Move data sufficiently so that no one erase block will wear out before the others
- Node types
 - inode: inode metadata and possibly a range of data. Data may be compressed
 - Directory entry: inode number, name, and version
 - Clear marker: newly-erased block may be safely used for storage

JFFS2 mounting

- Physical medium is scanned & checked for validity
 - Raw node references are allocated, inode cache structures allocated
- First pass through physical nodes
 - Build full map of data for each inode (and detect obsolete nodes)
- Second pass
 - Find inodes that have no remaining links on the file system and delete them
- Third pass
 - Free temporary info that was cached for each node

The End