

Operating Systems Design 17. Networking: Remote Procedure Calls

Paul Krzyzanowski
pxk@cs.rutgers.edu

1

Problems with the sockets API

The **sockets** interface forces a read/write mechanism

Programming is easier with a functional interface

2

RPC

1984: Birrell & Nelson

- Mechanism to call procedures on other machines

Remote Procedure Call

Goal: it should appear to the programmer that
a normal call is taking place

3

Regular procedure calls

You write:

```
x = f(a, "test", 5);
```

The compiler parses this and generates code to:

- Push the value 5 on the stack
- Push the address of the string "test" on the stack
- Push the current value of a on the stack
- Generate a call to the function f

In compiling f, the compiler generates code to:

- Push registers that will be clobbered on the stack to save the values
- Adjust the stack to make room for local and temporary variables
- Before a return, unadjust the stack, put the return data in a register, and issue a return instruction

4

Implementing RPC

No architectural support for remote procedure calls

Simulate it with tools we have
(local procedure calls)

Simulation makes RPC a
language-level construct

instead of an
operating system construct

The compiler
creates code to
send messages to
invoke remote
functions

The OS gives us
sockets

5

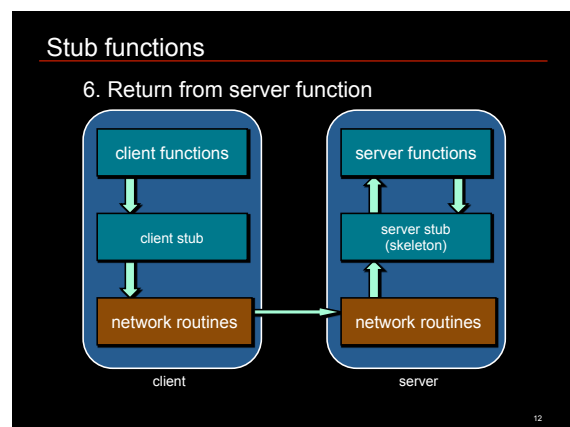
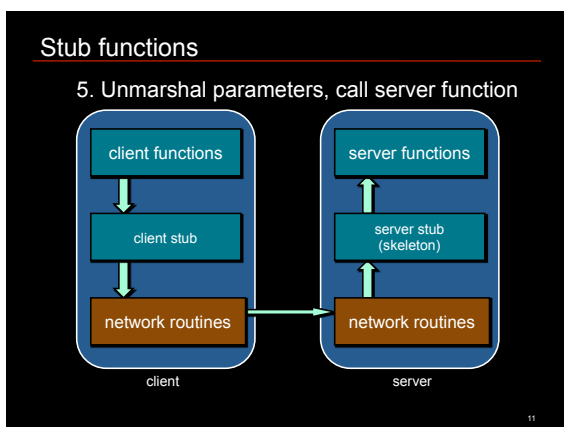
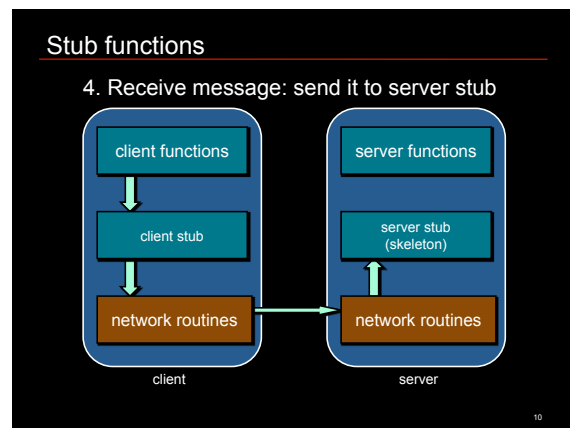
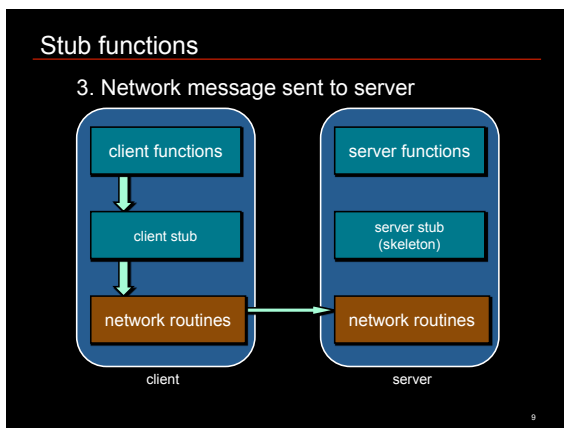
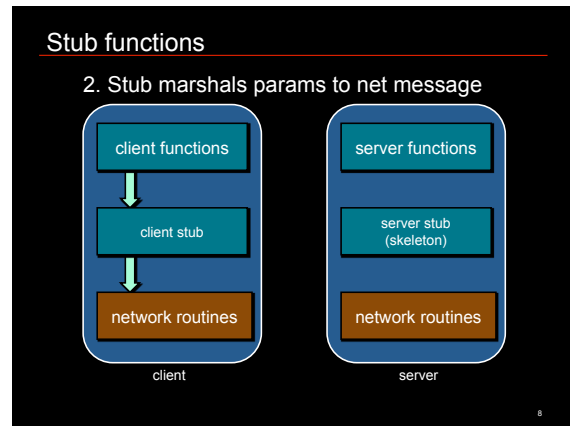
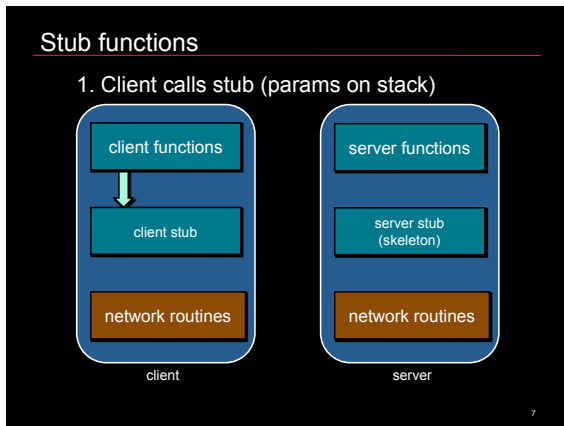
Implementing RPC

The trick:

Create **stub functions** to make it appear to the user that the call is
local


Stub function contains the function's interface

6



Stub functions


7. Marshal return value and send message



13

Stub functions


8. Transfer message over network



14

Stub functions


9. Receive message: client stub is receiver



15

Stub functions

10. Unmarshal return value, return to client code



16

Benefits

- Procedure call interface
- Writing applications is simplified
 - RPC hides all network code into stub functions
 - Application programmers don't have to worry about details
 - Sockets, port numbers, byte ordering

17

RPC has issues

18

Parameter passing

Pass by value

- Easy: just copy data to network message

Pass by reference

- Makes no sense without shared memory

19

Pass by reference?

1. Copy items referenced to message buffer
2. Ship them over
3. Unmarshal data at server
4. Pass *local* pointer to server stub function
5. Send new values back

To support complex structures

- Copy structure into pointerless representation
- Transmit
- Reconstruct structure with local pointers on server

20

Representing data

No such thing as
incompatibility problems on local system

Remote machine may have:

- Different byte ordering
- Different sizes of integers and other types
- Different floating point representations
- Different character sets
- Alignment requirements

21

Representing data

IP (headers) forced all to use **big endian** byte ordering for 16- and 32-bit values

Big endian: Most significant byte in low memory

- SPARC v9, Motorola 680x0, older PowerPC

Little endian: Most significant byte in high memory

- Intel IA-32, x64

Bi-endian: Processor may operate in either mode

- ARM, PowerPC, MIPS, SPARC V9, IA-64 (Intel Itanium)

```
main() {
    unsigned int n;
    char *a = (char *) &n;

    n = 0x11223344;
    printf("%02x, %02x, %02x, %02x\n",
           a[0], a[1], a[2], a[3]);
}
```

Output on an Intel:
44, 33, 22, 11

Output on a PowerPC:
11, 22, 33, 44

22

Representing data

Need standard encoding to enable communication between heterogeneous systems

- e.g. Sun's RPC uses XDR (eXternal Data Representation)
- ASN.1 (ISO Abstract Syntax Notation)
- JSON (JavaScript Object Notation)
- Google Protocol Buffers
- W3C XML Schema Language

23

Representing data

Implicit typing

- only values are transmitted, not data types or parameter info
- e.g., Sun XDR

Explicit typing

- Type is transmitted with each value
- e.g., ISO's ASN.1, XML

24

Where to bind?

Need to locate host and correct server process

25

Where to bind? – Solution 1

Maintain centralized DB that can locate a host that provides a particular service
(Birrell & Nelson's 1984 proposal)

26

Where to bind? – Solution 2

A server on each host maintains a DB of *locally* provided services

Solution 1 is problematic for Sun NFS – identical file servers serve different file systems

27

Transport protocol

TCP or UDP? Which one should we use?

- Some implementations may offer only one (e.g. TCP)
- Most support several
 - Allow programmer (or end user) to choose at runtime

28

When things go wrong

- Local procedure calls do not fail
 - If they core dump, entire process dies
- More opportunities for error with RPC
- Transparency breaks here
 - Applications should be prepared to deal with RPC failure

29

More issues

Performance

- RPC is slower ... a lot slower

Security

- messages visible over network
- Authenticate client
- Authenticate server

30

Programming with RPC

Language support

- Most programming languages have no concept of remote procedure calls
- Language compilers will not generate client and server stubs

Common solution:

- Use a separate compiler to generate stubs (pre-compiler)

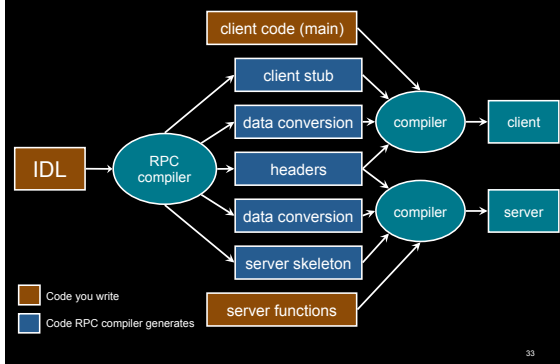
31

Interface Definition Language

- Allow programmer to specify remote procedure interfaces (names, parameters, return values)
- Pre-compiler can use this to generate client and server stubs:
 - Marshaling code
 - Unmarshaling code
 - Network transport routines
 - Conform to defined interface
- Similar to function prototypes

32

RPC compiler



33

The End

34