

Operating Systems Design

19. Protection

Paul Krzyzanowski
pxk@cs.rutgers.edu

Protection & Security

- Security
 - Prevention of unauthorized access to a system
 - Malicious or accidental access
 - “access” may be:
 - user login, a process accessing things it shouldn’t, physical access
 - The access operations may be reading, destruction, or alteration
- Protection
 - The mechanism that provides and enforces controlled access of resources to processes
 - A protection mechanism *enforces* security policies

Principle of Least Privilege

- At each abstraction layer, every element (user, process, function) should be able to access *only* the resources necessary to perform its task
- Even if an element is compromised, the scope of damage is limited
- Consider:
 - Violation: a compromised print daemon allows one to add users
 - Violation: a process can write a file even though there is no need to
 - Violation: admin privileges set by default for any user account
 - Good: Private member functions & Local variables in functions limit scope
- Least privilege is often difficult to define & enforce

Privilege Separation

- Divide a program into multiple parts: high & low privilege components
- Example on POSIX systems
 - Each process has a *real* and *effective* user ID
 - Privileges are evaluated based on the effective user ID
 - Normally, *uid* == *euid*
 - An executable file may be tagged with a *setuid bit*
 - `chmod +sx filename`
 - When run, *uid* = user's ID; *euid* = file owner's ID
 - Separating a program
 - Run a *setuid* program
 - Create a communication link to self (*pipe*, *socket*, shared memory)
 - *fork*
 - One process will call *seteuid(getuid())* to lower its privilege

Security Goals

- **Authentication**
 - Ensure that users, machines, programs, and resources are properly identified
- **Confidentiality**
 - Prevent unauthorized access to data
- **Integrity**
 - Verify that data has not been compromised: deleted, modified, added
- **Availability**
 - Ensure that the system is accessible

The Operating System

- The OS provides processes with access to resources

Resource	OS component
Processor(s)	Process scheduler
Memory	Memory Management + MMU
Peripheral devices	Device drivers & buffer cache
Logical persistent data	File systems
Communication networks	Sockets

- Resource access attempts go through the OS
- OS decides whether access should be granted
 - Rules that guide the decision = *policy*

Domains of protection

- Processes interact with objects
 - Objects:
 - hardware (CPU, memory, I/O devices)
 - software: files, semaphores, messages, signals
- A process should be allowed to access only objects that it is authorized to access
 - A process operates in a **protection domain**
 - Protection domain defines the objects the process may access and how it may access them

Modeling Protection: Access Matrix

- Rows: domains
- Columns: objects
- Each entry represents an access right of a domain on an object

		<i>objects</i>		
<i>domains of protection</i>		F_0	F_1	<i>Printer</i>
	D_0	read	read-write	print
	D_1	read-write-execute	read	
	D_2	read-execute		
	D_3		read	print
	D_4			print

Access Matrix: Domain Transfers

- Switching from one domain to another is a configurable policy

domains of protection	objects							
	F_0	F_1	Printer	D_0	D_1	D_2	D_3	D_4
	D_0	read	read-write	print	–	switch	switch	
	D_1	read-write-execute	read		–			
	D_2	read-execute			switch	–		
	D_3		read	print				
	D_4			print				

Access Matrix: Additional operations

- **Copy:** allow delegation of rights
 - Copy a specific access right on an object from one domain to another
 - Rights may specify either a copy or a transfer of rights

domains of protection

	F_0	F_1	<i>Printer</i>	D_0	D_1	D_2	D_3	D_4
D_0	read	read-write	print	–	switch	sv		
D_1	read-write-execute	read*						
D_2	read-execute				switch	–		
D_3		read	print					
D_4			print					

E.g., a process executing in D_1 can give a read right on F_1 to domain D_2

Access Matrix: Additional operations

- **Owner**: allow new rights to be added or removed
 - An object may be identified as being *owned* by the domain
 - Owner can add and remove any right in any column of the object

objects

	F_0	F_1	<i>Printer</i>	D_0	D_1	D_2	D_3	D_4
<i>domains of protection</i>	D_0	read owner	read-write	print	–	switch	sw	
	D_1	read-write-execute	read*		–			
	D_2	read-execute			switch			
	D_3		read	print				
	D_4		print					

E.g., a process executing in D_0 can give a read right on F_0 to domain D_3 and remove the execute right from D_1

Access Matrix: Additional operations

- **Control**: change entries in a row
 - If $\text{access}(i, j)$ includes a *control right*, then a process executing in Domain i can change access rights for Domain j

objects

	F_0	F_1	<i>Printer</i>	D_0	D_1	D_2	D_3	D_4
<i>domains of protection</i>	D_0	read owner	read-write	print	–	switch	switch	
	D_1	read-write-execute	read*		–			control
	D_2	read-execute				switch		
	D_3		read	print				
	D_4			print				

E.g., a process executing in D_1 can modify any rights in domain D_4

Implementing an access matrix

- A single table is usually impractical
 - Big size: $\# \text{ domains (users)} \times \# \text{ objects (files)}$
 - Objects may come and go frequently
- Access Control List
 - Associate a column of the table with each object

Implementing an access matrix

- **Access Control List**
 - Associate a column of the table with each object

domains of protection

	<i>F_0</i>	<i>F_1</i>	<i>Printer</i>	<i>D_0</i>	<i>D_1</i>	<i>D_2</i>	<i>D_3</i>	<i>D_4</i>
<i>D_0</i>	read owner	read-write	print					
<i>D_1</i>	read-write-execute	read*			–			
<i>D_2</i>	read-execute				switch	–		
<i>D_3</i>		read	print					
<i>D_4</i>			print					

objects

ACL for file F_0

Example: Limited ACLs in POSIX systems

- Problem: an ACL takes up a varying amount of space (possibly a lot!)
 - Won't fit in an inode
- UNIX Compromise:
 - A file defines access rights for three domains: the owner, the group, and everyone else
 - Permissions
 - Read, write, execute, directory search
 - Set user ID on execution
 - Set group ID on execution
 - Default permissions set by the *umask* system call
 - *chown* system call changes the object's owner
 - *chmod* system call changes the object's permissions

Example: Full ACLs in POSIX systems

- *What if we really want an ACL?*
- Extended attributes: stored outside of the inode
- Enumerated list of permissions on users and groups
 - Operations on all objects:
 - *delete, readattr, writeattr, readextattr, writeextattr, readsecurity, writesecurity, chown*
 - Operations on directories
 - *list, search, add_file, add_subdirectory, delete_child*
 - Operations on files
 - *read, write, append, execute*
 - Inheritance controls

Implementing an access matrix

- **Capability List**

- Associate a row of the table with each domain

objects

domains of protection

	F_0	F_1	<i>Printer</i>	D_0	D_1	D_2	D_3	D_4
D_0	read owner	read-write	print	–	switch	swtich		
D_1	read-write-execute	read*			–			
D_2	read-execute				switch	–		
D_3		read	print					
D_4			print					

Capability list for domain D_1

Capability Lists


- List of objects together with the operations allowed on the objects
- Each item in the list is a *capability*: the operations allowed on a specific object
- A process presents the capability along with a request
 - Possessing the capability means that access is allowed
- A process cannot modify its capability list

Access Control Models: MAC vs. DAC

- **DAC: Discretionary Access Control**
 - A subject (domain) can pass information onto any other subject
 - In some cases, access rights may be transferred
 - *Most systems use this*
- **MAC: Mandatory Access Control**
 - Policy is centrally controlled
 - Users cannot override the policy

Multi-level Access Control

- Typical MAC implementations use a **Multi-Level Secure (MLS)** access model
- **Bell-LaPadula** model
 - Identifies the ability to access and communicate data
 - Objects are classified into a hierarchy of sensitivity levels
 - Unclassified, Confidential, Secret, Top Secret
 - Users are assigned a clearance
 - “**No read up; no write down**”
 - Cannot read from a higher clearance level
 - Cannot write to a lower clearance level
- Works well for government information
- Does not translate well to civilian life

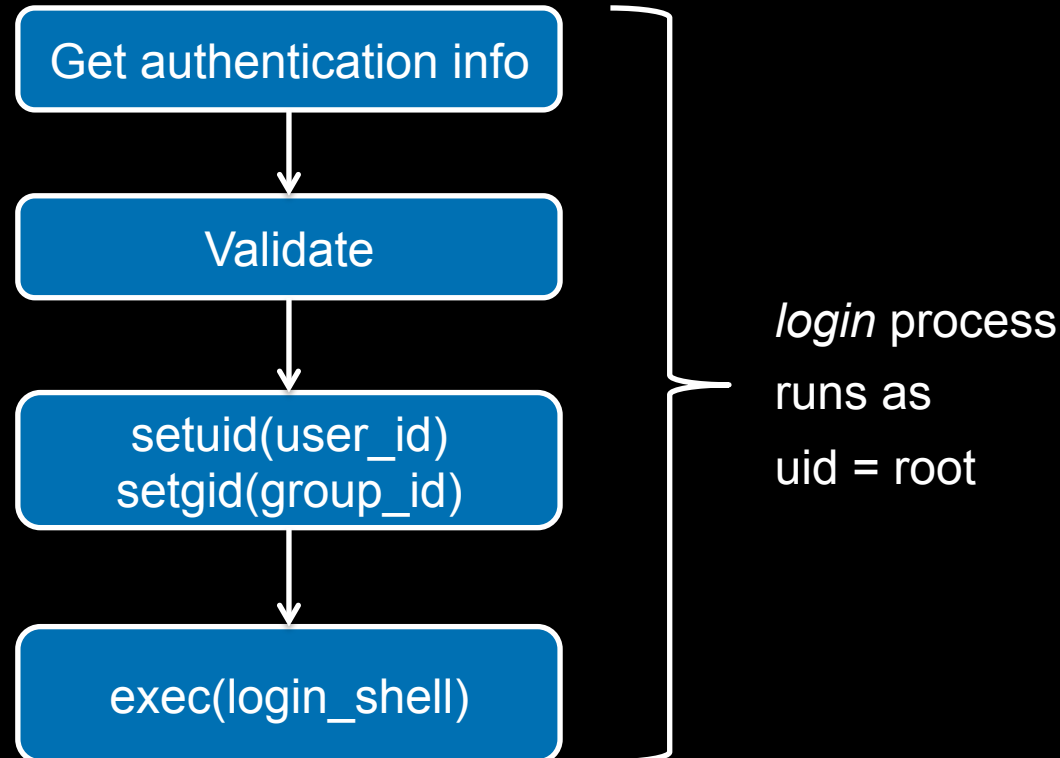


*Confidential cannot read Secret
Confidential cannot write Unclassified*

Authentication

Authentication

- Establish & verify identity
 - Then decide whether to allow access to resources
- Example: login



Password Authentication Protocol (PAP)

- Reusable passwords
- Server keeps a database of *username:password* mappings
- Prompt client/user for a login name & password
- To authenticate, use the login name as a key to look up the corresponding password in a database (file) to authenticate

```
if (supplied_password == retrieved_password)
    then user is authenticated
```

PAP: Reusable passwords

One problem: what if the password file isn't sufficiently protected and an intruder gets hold of it, he gets all the passwords!

Enhancement:

Store a hash of the password in a file

- given a file, you don't get the passwords
- have to resort to a dictionary or brute-force attack
- Unix approach
 - Password encrypted with 3DES hashes; then MD5 hashes; now SHA512 hashes
 - Salt used to guard against **dictionary attacks**

Authentication

Three factors:

- something you have *key, card*
 - can be stolen
- something you know *passwords*
 - can be guessed, shared, stolen
- something you are *biometrics*
 - costly, can be copied (sometimes)

Authentication

factors may be combined

– ATM machine: 2-factor authentication

- **ATM card** *something you have*
- **PIN** *something you know*

Identification vs. Authentication

- **Identification:**
 - Who are you?
 - User name, account number, ...
- **Authentication:**
 - Prove it!
 - Password, PIN, encrypt nonce, ...

Versus Authorization

Authorization defines access control

Once we know a user's identity:

- Allow/disallow request
- Operating system enforces system access based on user's credentials
 - Network services usually run in another context
 - Network server may not know of the user
 - Application takes responsibility
- May contact an authorization server
 - Trusted third party that will grant credentials
 - Kerberos ticket granting service
 - **RADIUS** (centralized authentication/authorization)

Three (Four?) A's of Security

- Authentication
 - Validate an identity or a message
- Authorization (Access Control)
 - Enforce policy
- Accounting
- Auditing

Accounting

If security has been compromised

... *what happened?*

... *who did it?*

... *how did they do it?*

Log transactions

- Logins
- Commands
- Database operations
- *Who looks at audits?*

Log to remote systems

- Minimize chances for intruders to delete logs

Auditing

Go through software source code and search for security holes

- Need access to source
 - Some operating systems > 50 million lines!
- Experienced staff + time
- E.g., OpenBSD

Complex systems will have more bugs

- And will be harder to audit

The End